



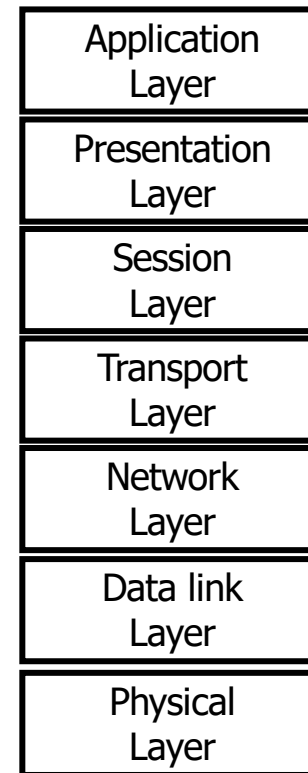
Computer Networks

Error Control

Chapter


1. Introduction
2. Protocols
3. Application layer
4. Web services
5. Distributed hash tables
6. Time synchronization
7. **Error control**
 - **Stop-and-Wait**
 - **Sliding window**
 - **Performance analysis**
8. Transport layer
9. TCP performance
10. Network layer
11. Internet protocol
12. Data link layer

Top-Down-Approach



Error control

Error control



- Noise, interference, buffer overflow, failures of components, etc. lead to bit errors and packet loss
- This can be compensated by protocols supporting error detection, acknowledgments, and repeated transmissions

Error control

- We study three fundamental protocols for error control
 - **Stop-and-Wait**
 - Sender adds checksum (best a cyclic redundancy check, CRC) to the packet in order to identify bit errors
 - Receiver sends acknowledgment (ACK) if packet was received w/o errors
 - If acknowledgement is not received until a deadline (timeout), the packet will be sent again
 - As duplicates can be generated, sequence numbers (SQN) are added to identify such duplicates
 - Bandwidth-delay-product has huge impact; if very large, the sender will spend most of the time waiting for acknowledgments
 - Sliding window protocols
 - Multiple packets are sent in a burst to fill the channel between sender and receiver
 - **Go-Back-N** and **Selective Repeat**
 - Protocols differ in terms of timeout, acknowledgment handling, and repeated transmissions

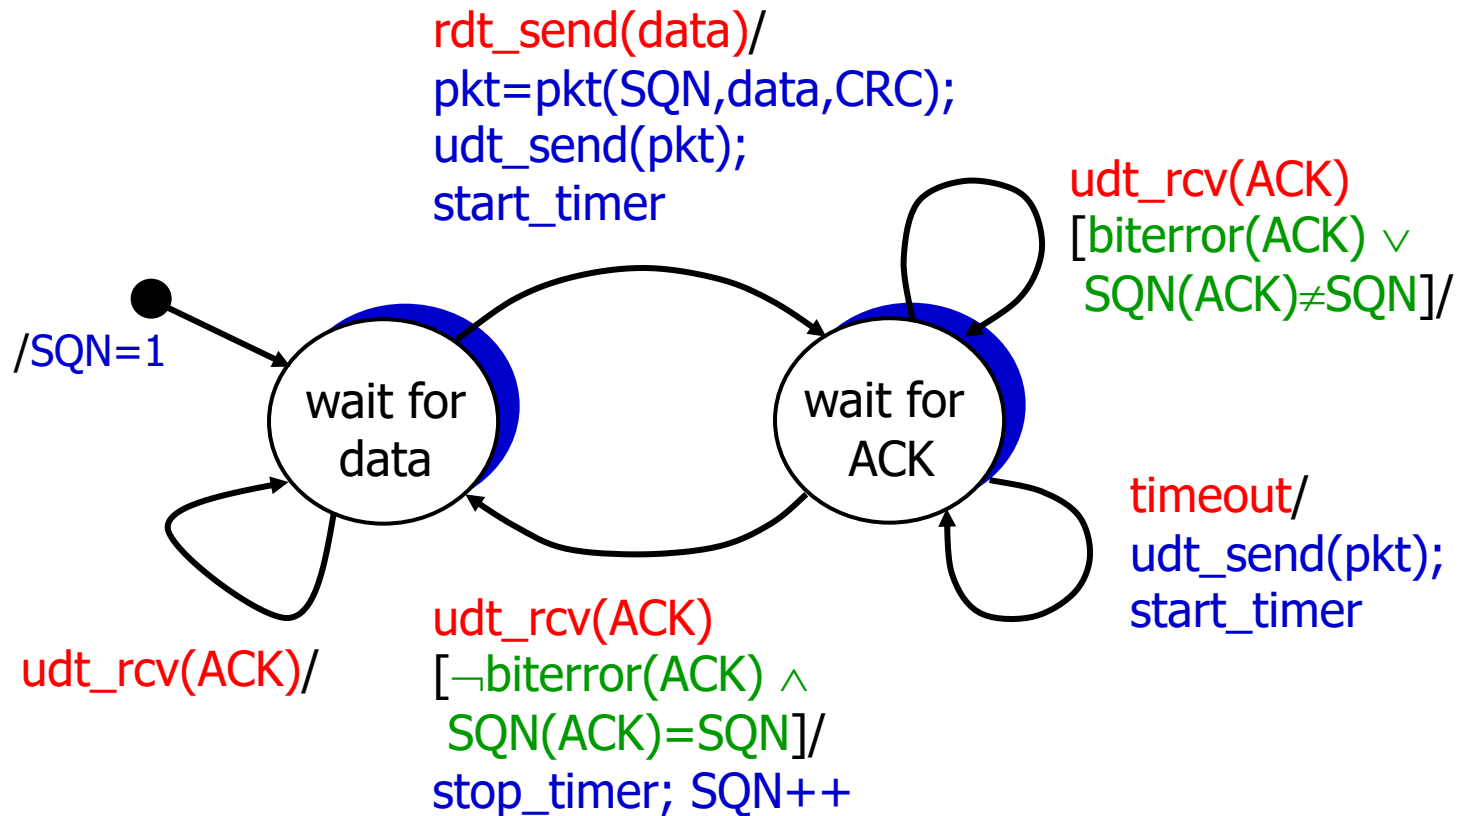
Stop-and-Wait

Stop-and-Wait

- How does stop-and-wait work?
- Informal:
 - Senders
 1. Send packet with current SQN, start timer
 2. If correct ACK is received, increment SQN and restart at step 1
 3. If timeout occurs, resend packet and restart timer
 - Receiver
 - If correct packet is received (no bit error, current SQN), send ACK with current SQN then increment SQN, otherwise send last ACK again

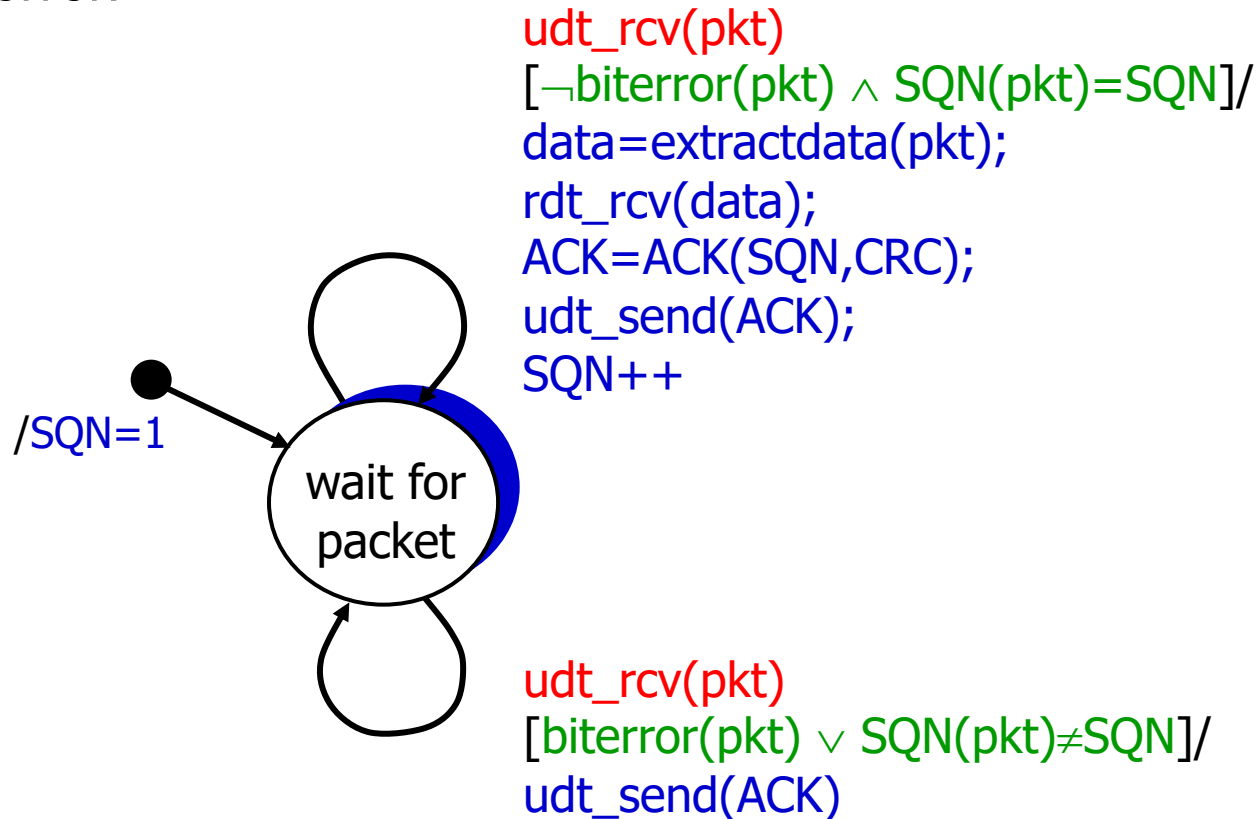
Stop-and-Wait

■ Sender:




Stop-and-Wait


■ Receiver:




Stop-and-Wait: normal execution




time




Stop-and-Wait: normal execution




time



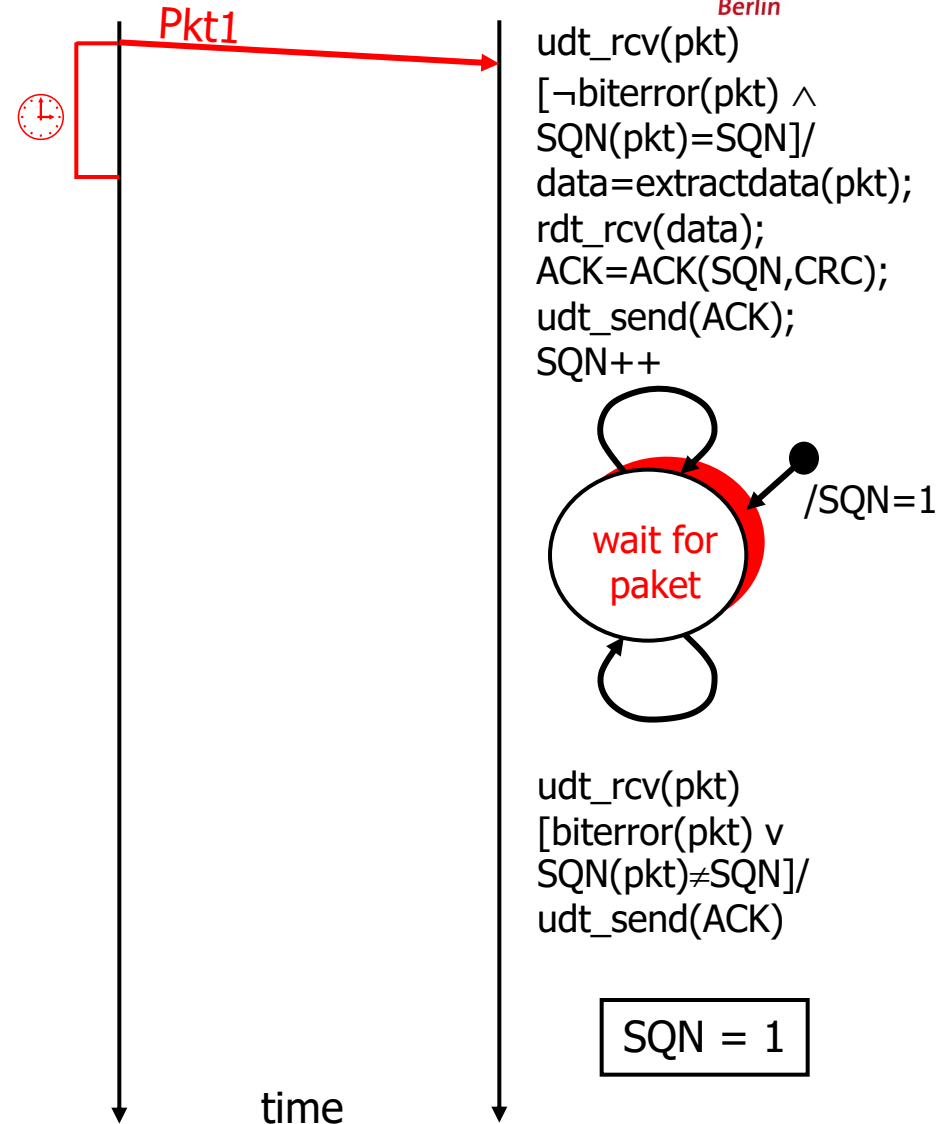

Stop-and-Wait: normal execution



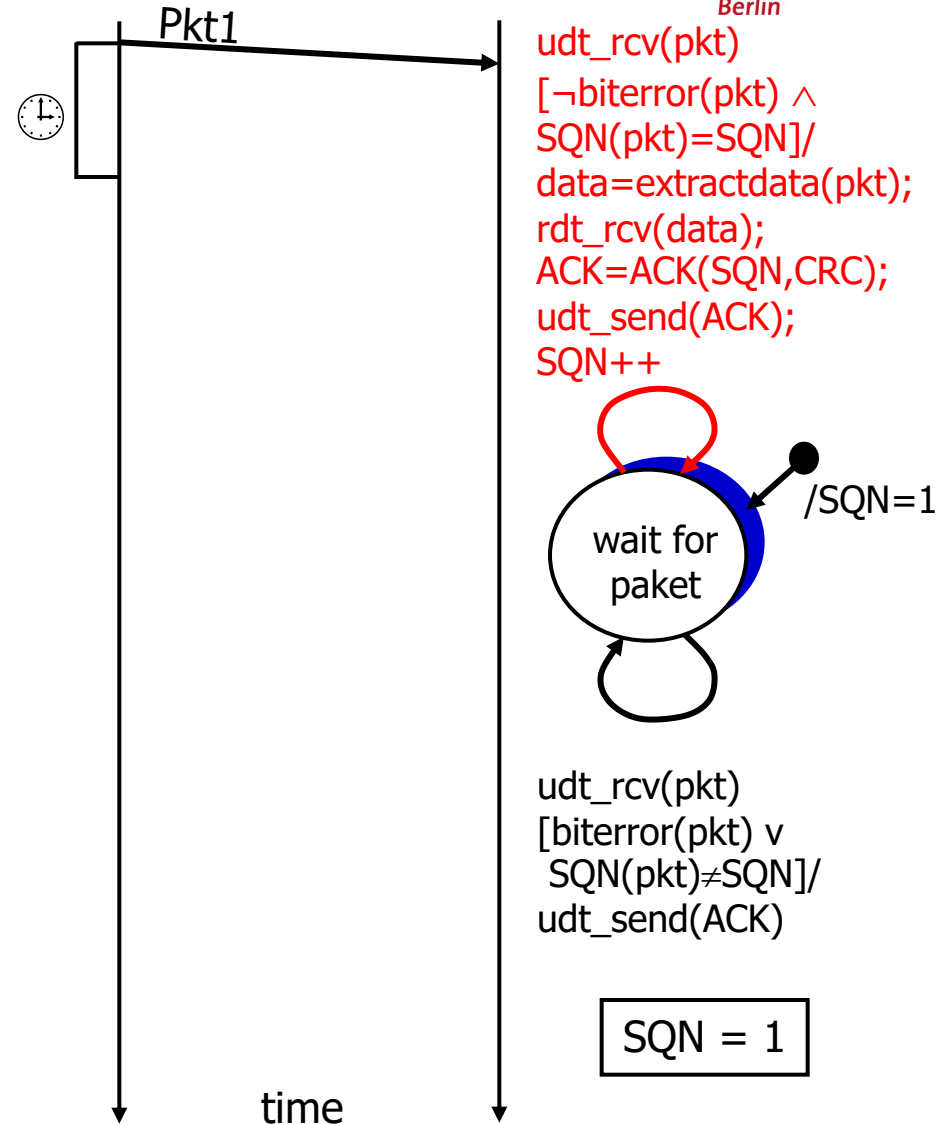

time



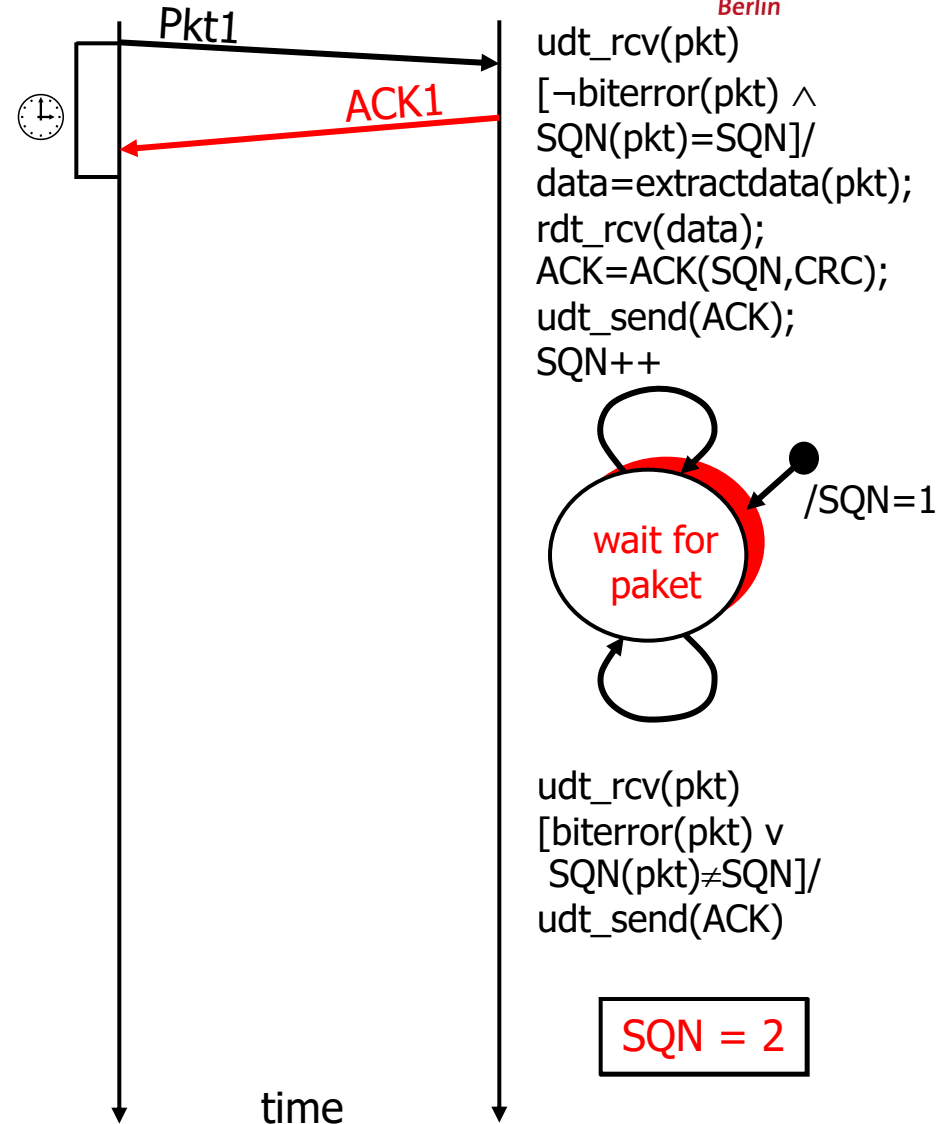

Stop-and-Wait: normal execution





Stop-and-Wait: normal execution





Stop-and-Wait: normal execution





Stop-and-Wait: normal execution





Stop-and-Wait: normal execution





Stop-and-Wait: packet loss





Stop-and-Wait: packet loss





Stop-and-Wait: packet loss




Stop-and-Wait: packet loss

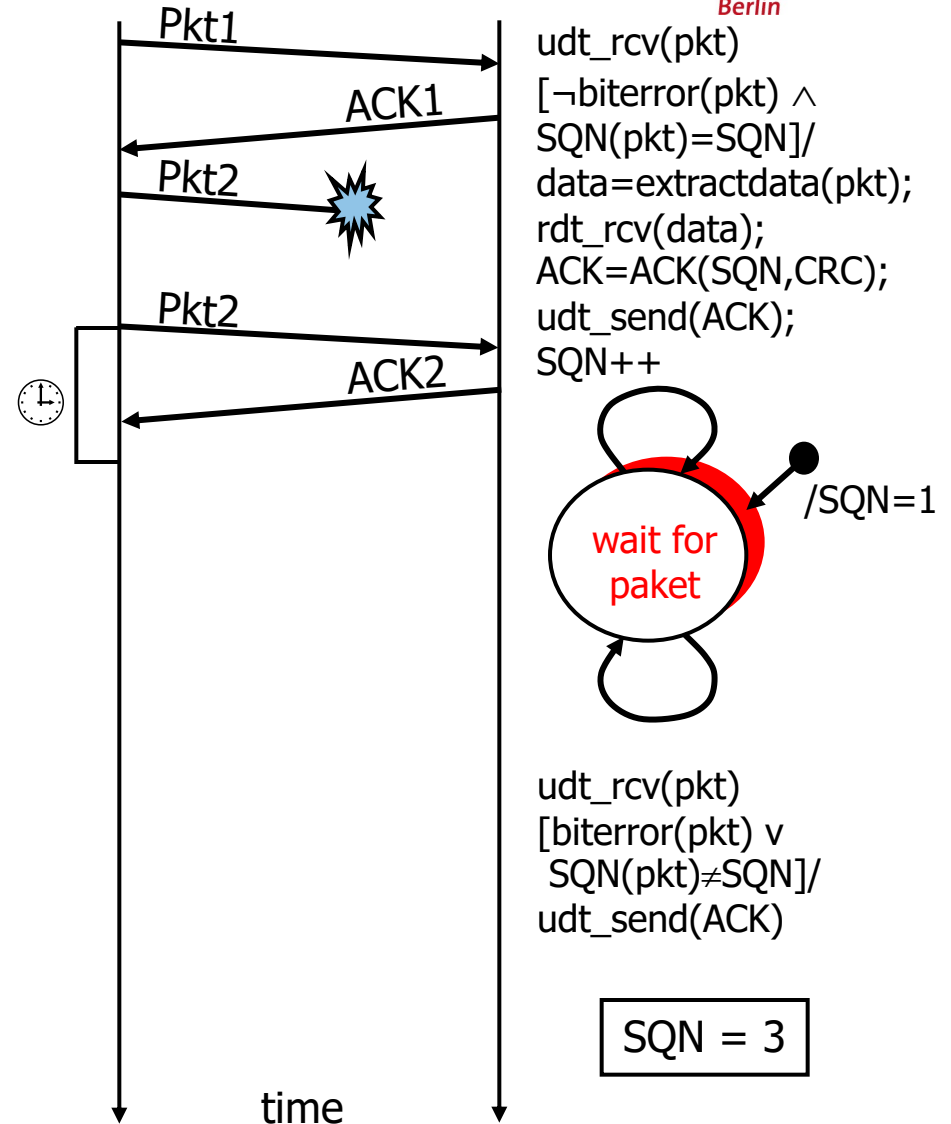



Stop-and-Wait: packet loss









Stop-and-Wait: packet loss

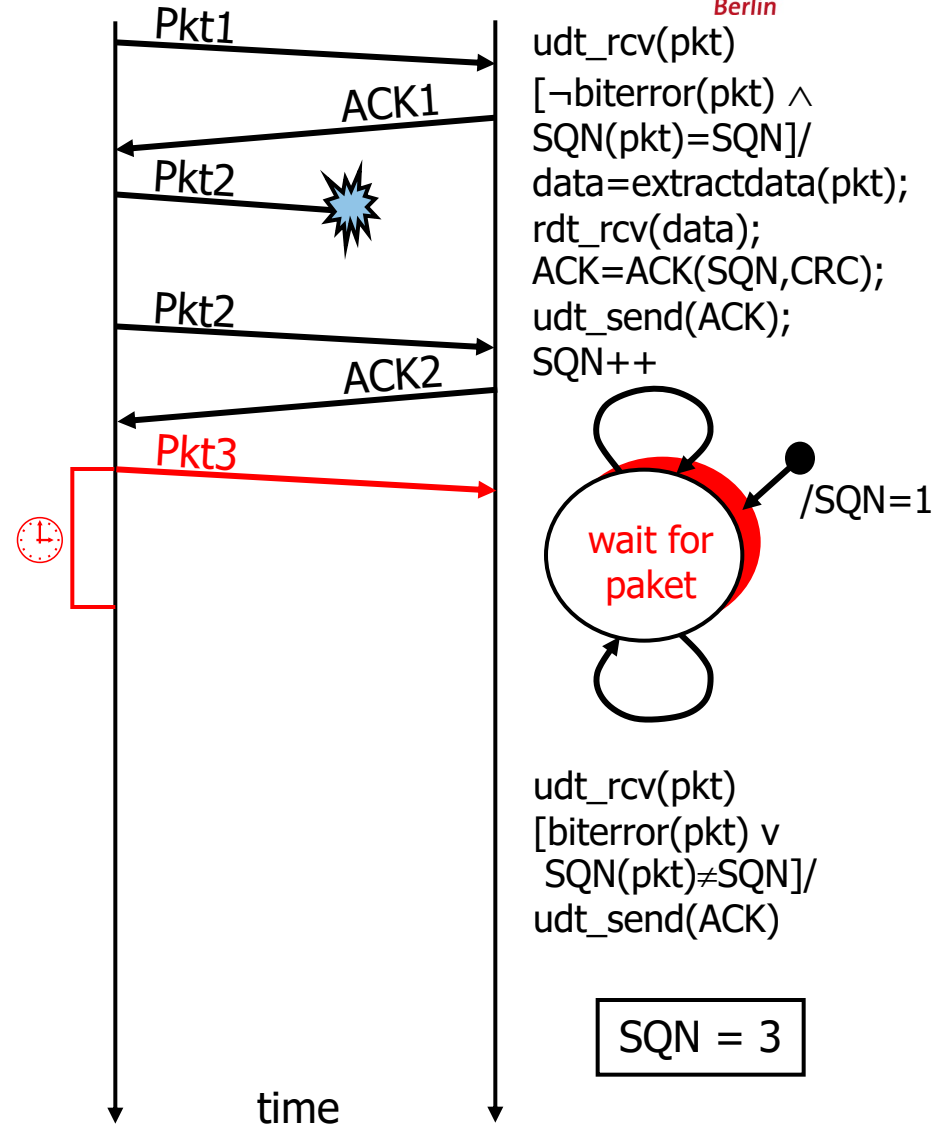


Stop-and-Wait: packet loss

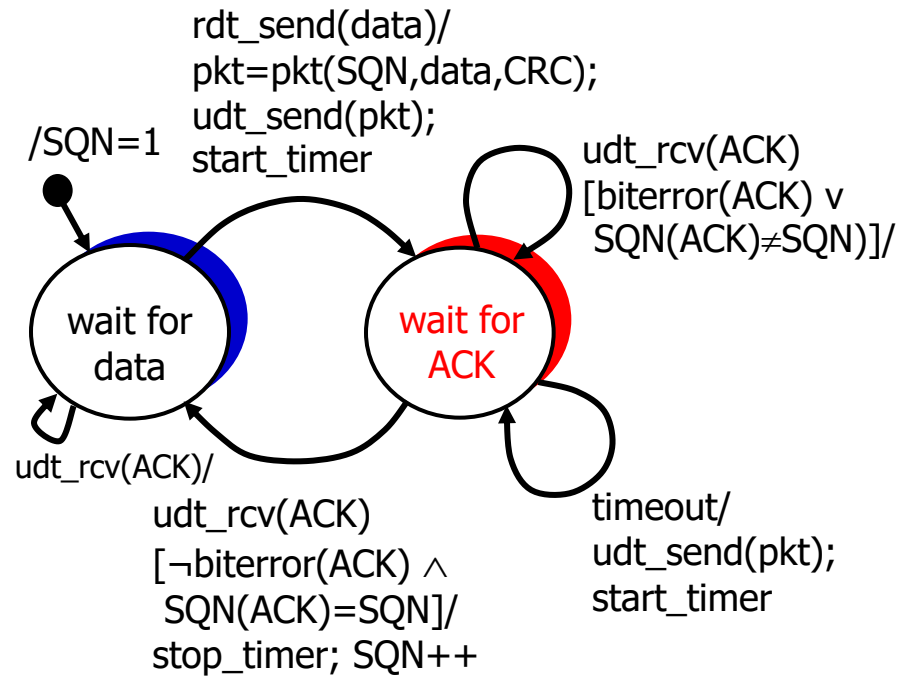


Stop-and-Wait: loss of ACK

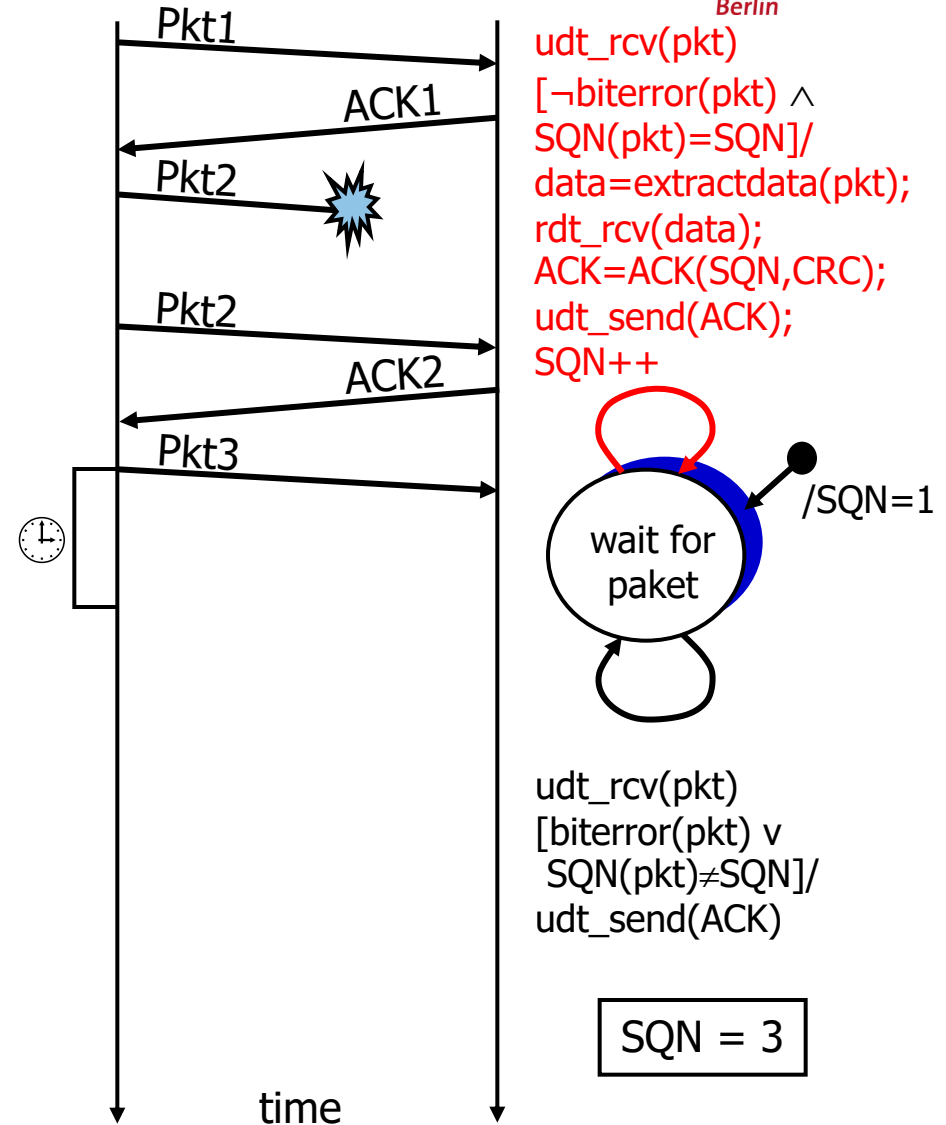




Stop-and-Wait: loss of ACK




SQN = 3




SQN = 3

Stop-and-Wait: loss of ACK

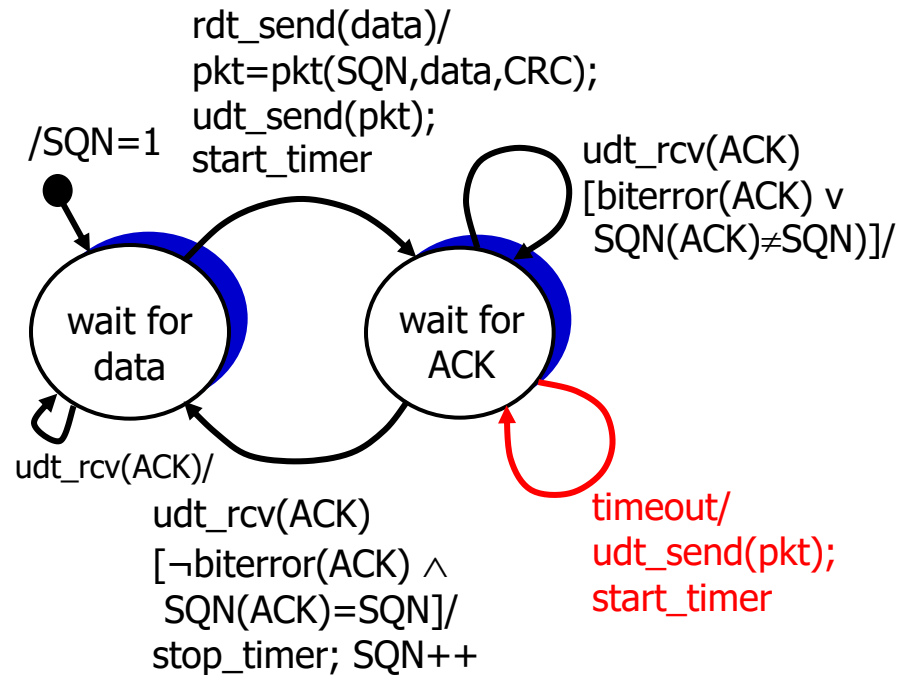


SQN = 3





SQN = 4

Stop-and-Wait: loss of ACK




SQN = 3




SQN = 4

Stop-and-Wait: loss of ACK




SQN = 3

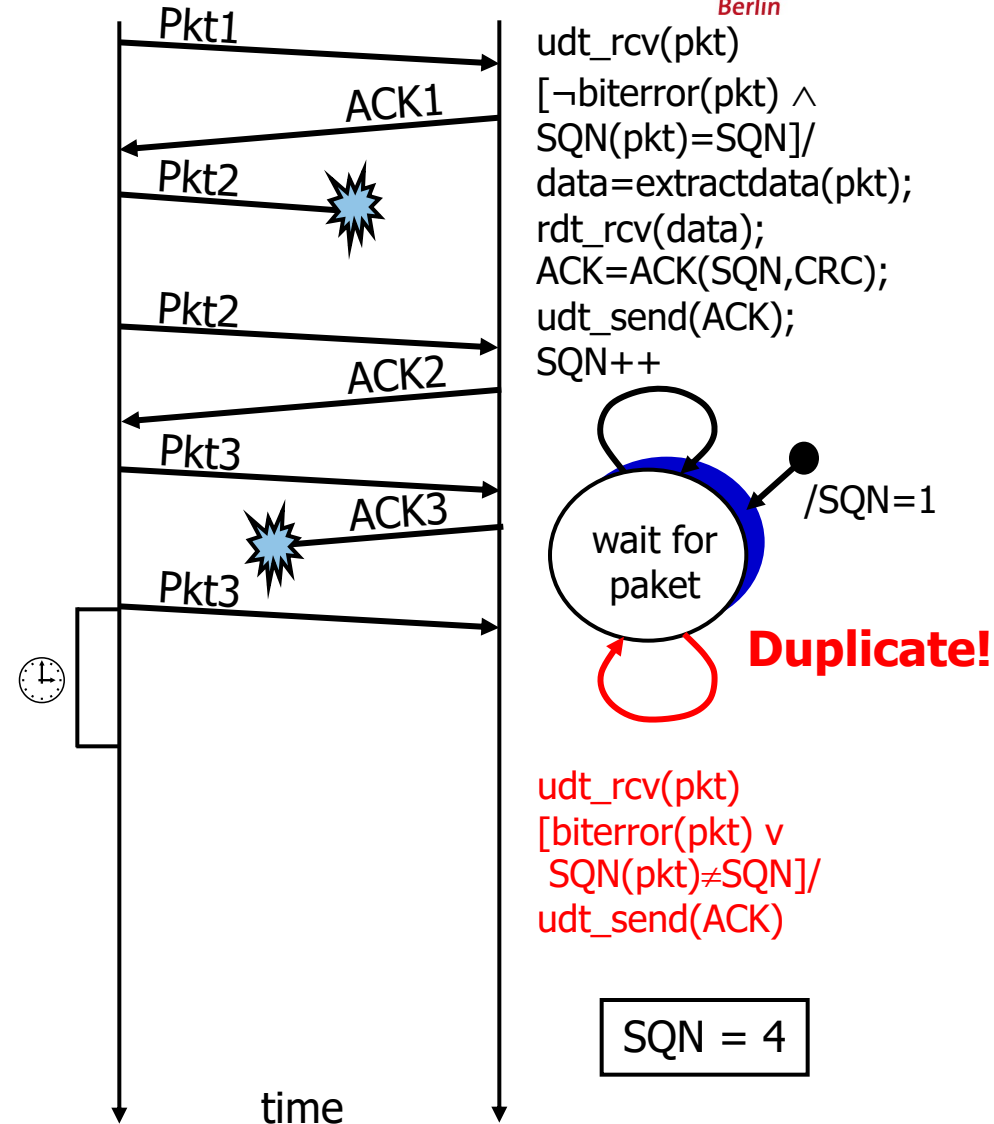


SQN = 4

Stop-and-Wait: loss of ACK




SQN = 3




SQN = 4

Stop-and-Wait: loss of ACK

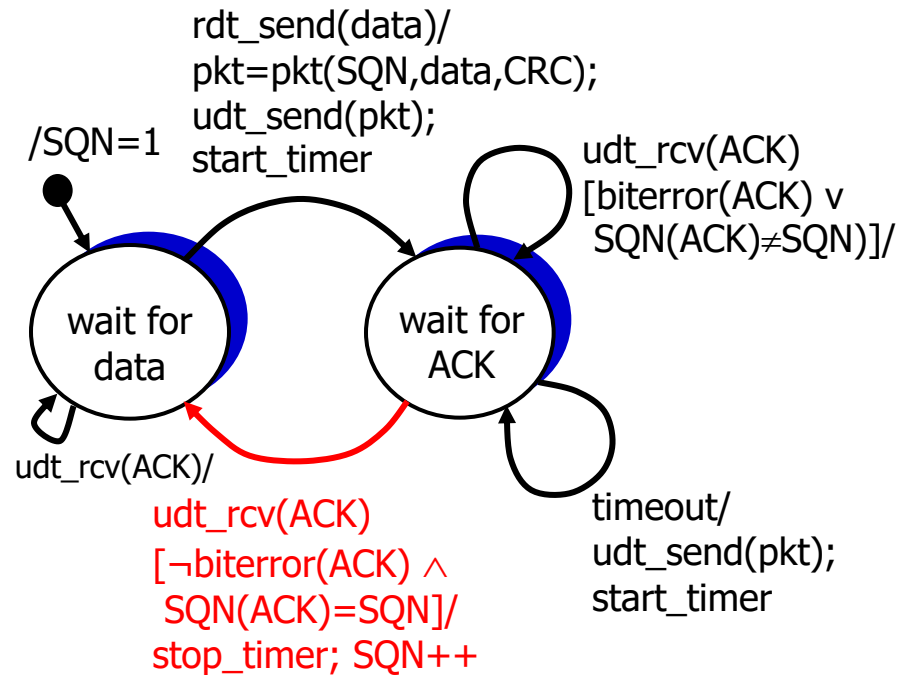


SQN = 3




SQN = 4

Stop-and-Wait: loss of ACK

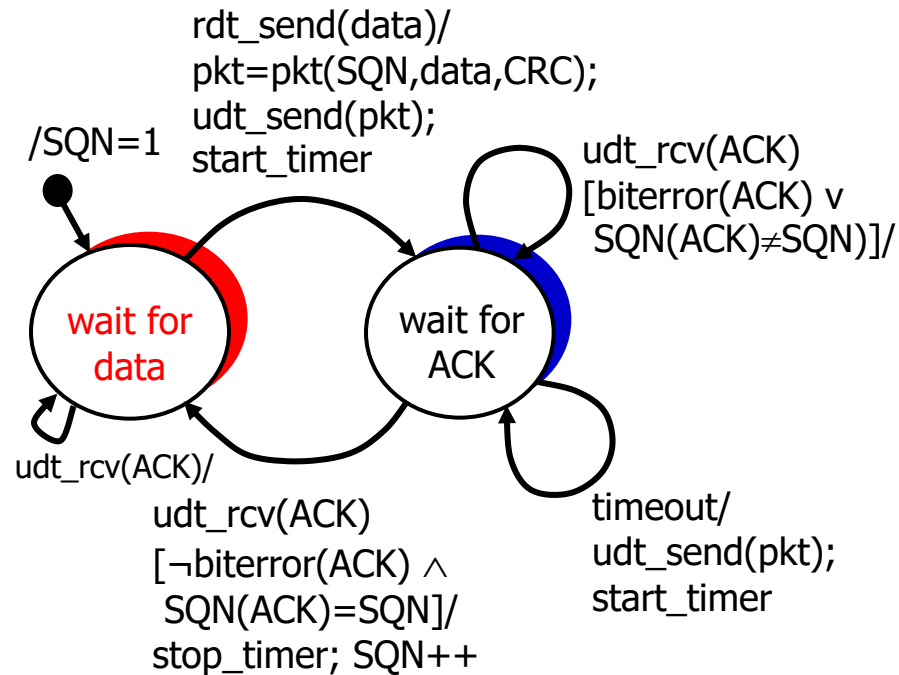


SQN = 3




SQN = 4

Stop-and-Wait: loss of ACK





SQN = 4




SQN = 4


Stop-and-Wait: delayed ACK



Stop-and-Wait: delayed ACK





SQN = 4





SQN = 4



Stop-and-Wait: delayed ACK





Stop-and-Wait: delayed ACK



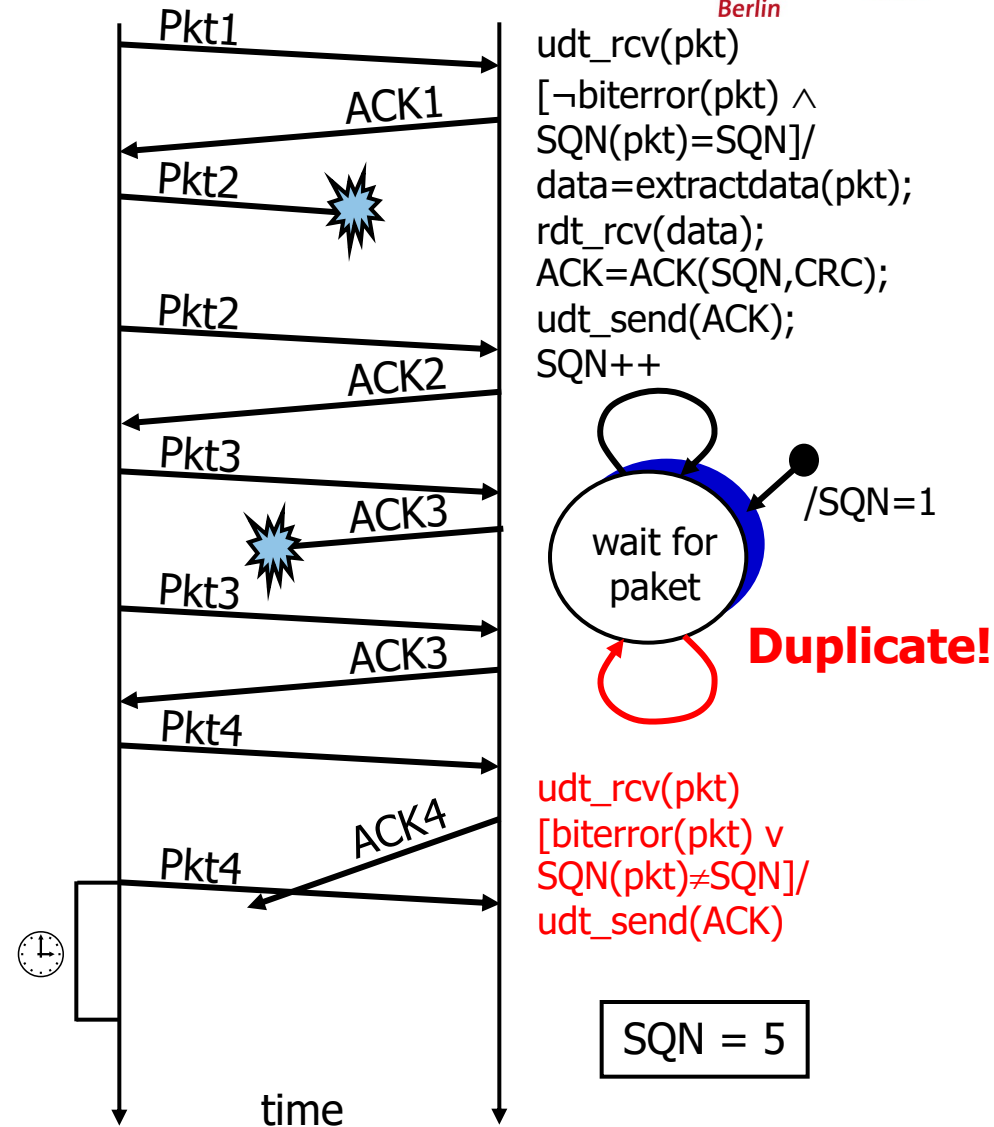

Stop-and-Wait: delayed ACK





Stop-and-Wait: delayed ACK




Stop-and-Wait: delayed ACK




Stop-and-Wait: delayed ACK



Stop-and-Wait: delayed ACK





SQN = 4





SQN = 5

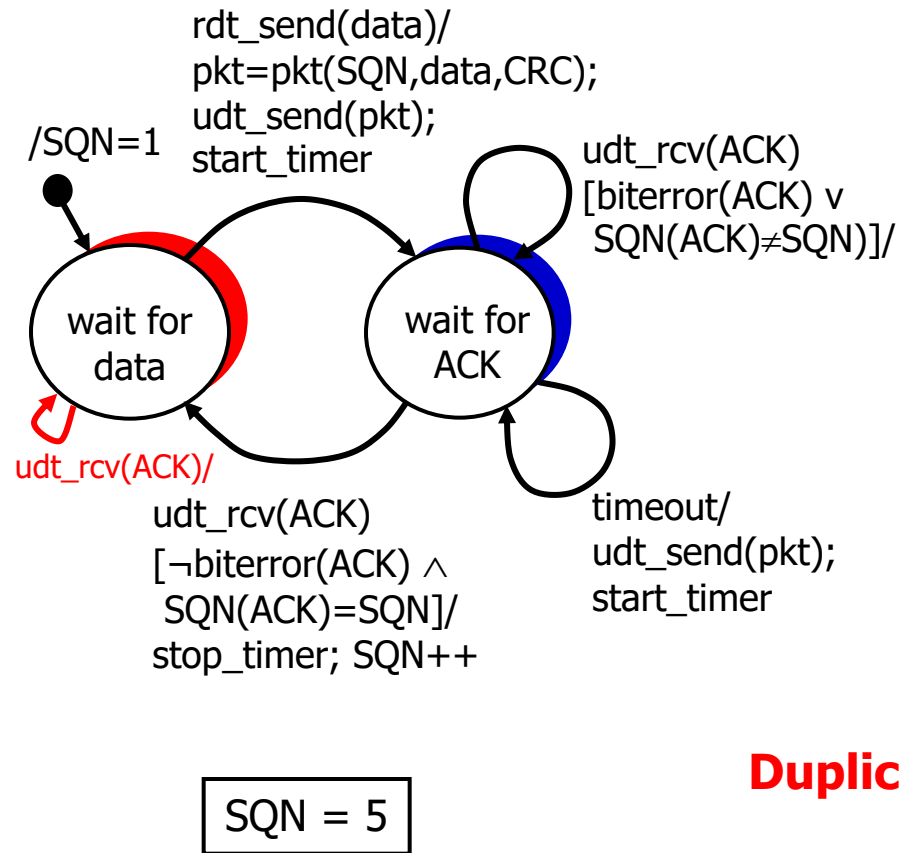
Stop-and-Wait: delayed ACK




Stop-and-Wait: delayed ACK



Stop-and-Wait: delayed ACK



Duplicate!



Stop-and-Wait



■ Sequence numbers

- Every implementation uses finite sequence numbers: a variable of n bit allows the use of 2^n sequence numbers
- Cyclic repeated use
- For Stop-and-Wait, one bit is sufficient to represent two sequence numbers (0 and 1)
- Stop-and-Wait using 0 and 1 as sequence numbers is also known as **alternating-bit-protocol**

Sliding window protocols

Go-Back-N and Selective Repeat

- Stop-and-Wait is inefficient for large bandwidth-delay-products
- **Sliding window protocols** send multiple packets before requiring an ACK




Go-Back-N

■ Principles

- Sender can send multiple packets (up to some maximum) before receiving the first ACK
- A timer is started after the transmission of the first packet
- All non-acknowledged packets are buffered
- After timeout, **all unacknowledged packets are sent again**
- The receiver sends **cumulative ACKs**: each ACK acknowledges all packets up to and including the current SQN
- The receiver accepts packets only in the correct order and **does not need a buffer**

Go-Back-N

■ Send buffer



- base: SQN of the oldest unacknowledged packet
- nextSQN: SQN of the next packet to be transmitted
- W: window size, number of packets that can be sent before ACK
- Window $[base, base+W-1]$ is constantly moved from left to right; thanks to cumulative ACKS, the structure remains unchanged:
 - $[base, nextSQN-1]$: packets that have been sent but not yet acknowledged
 - $[nextSQN, base+W-1]$: packets that have not yet been sent but can be sent before next ACK

Go-Back-N

■ Informal description


■ Sender

1. If data available and the window is not full: send packet with nextSQN, increment nextSQN, if this was the first packet in the window, start timer
2. If ACK is in window, move window to this SQN; if the window is now empty, stop timer; otherwise, re-start timer
3. If timeout, send all packets in window again, restart timer

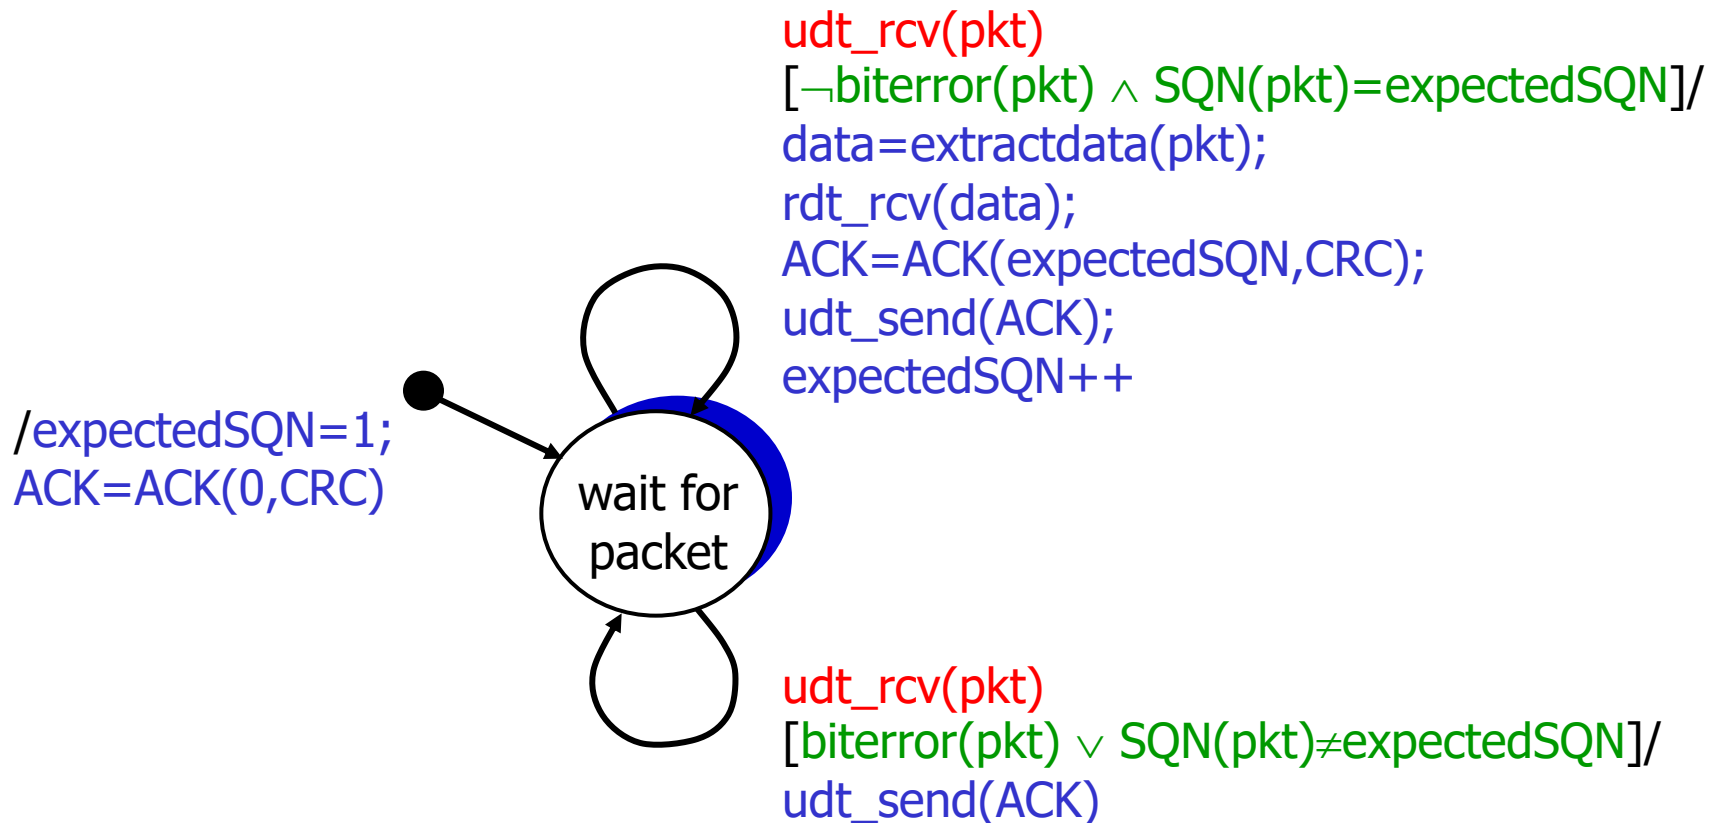
■ Receiver

- If packet is correct contains next expected SQN, send ACK with current SQN, increment SQN; otherwise re-send last ACK

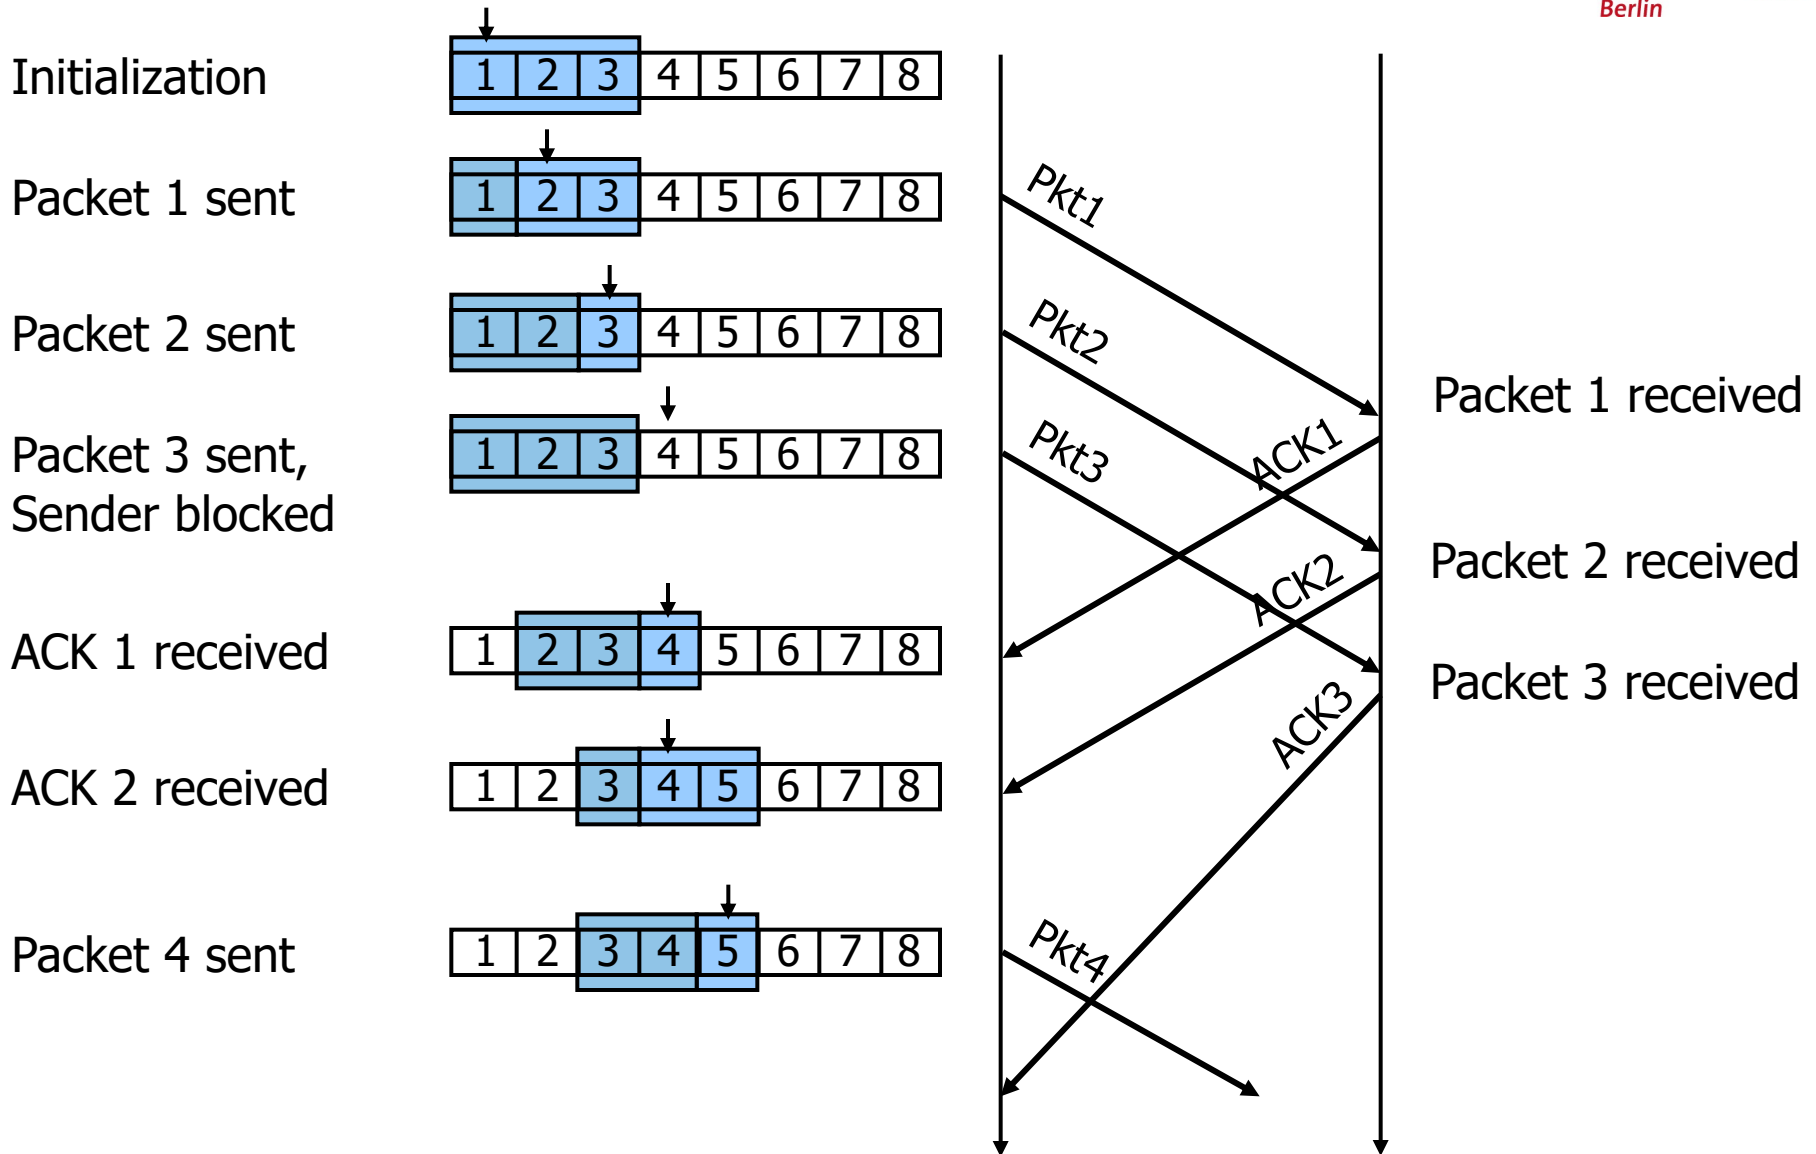
Go-Back-N: Sender




Go-Back-N: Receiver




Go-Back-N: normal execution





Go-Back-N: packet loss



Timer neu starten



Timeout, all not yet acknowledged packets are re-transmitted




Packet 1 received

Packet 2 received



Packet 3 received

Packet 4 received

Go-Back-N: lost and delayed ACKs



cumulative ACK
to compensate
loss and delay





Selective Repeat

■ Principles

- Sender can send multiple packets (up to some maximum) before receiving the first ACK
- A timer is started on a per-transmitted-packet-basis
- All non-acknowledged packets are buffered
- After timeout, **the associated packet is re-sent**
- Receiver sends **selective ACKs**: the semantics is that the packet with the SQN in the ACK was successfully received
- Receiver needs a **buffer** to compensate for (still) missing packets

Selective Repeat: buffer management



- base, nextSQN, W (as for Go-Back-N)
- Window at sender contains **sent but not yet acknowledged**, **sent and acknowledged**, and **unsent packets**
- Receiver buffers all received packets
- Window at receiver contains **received packets**, **missing packets**, and space for **not yet received packets**

Selective Repeat

■ Informal description

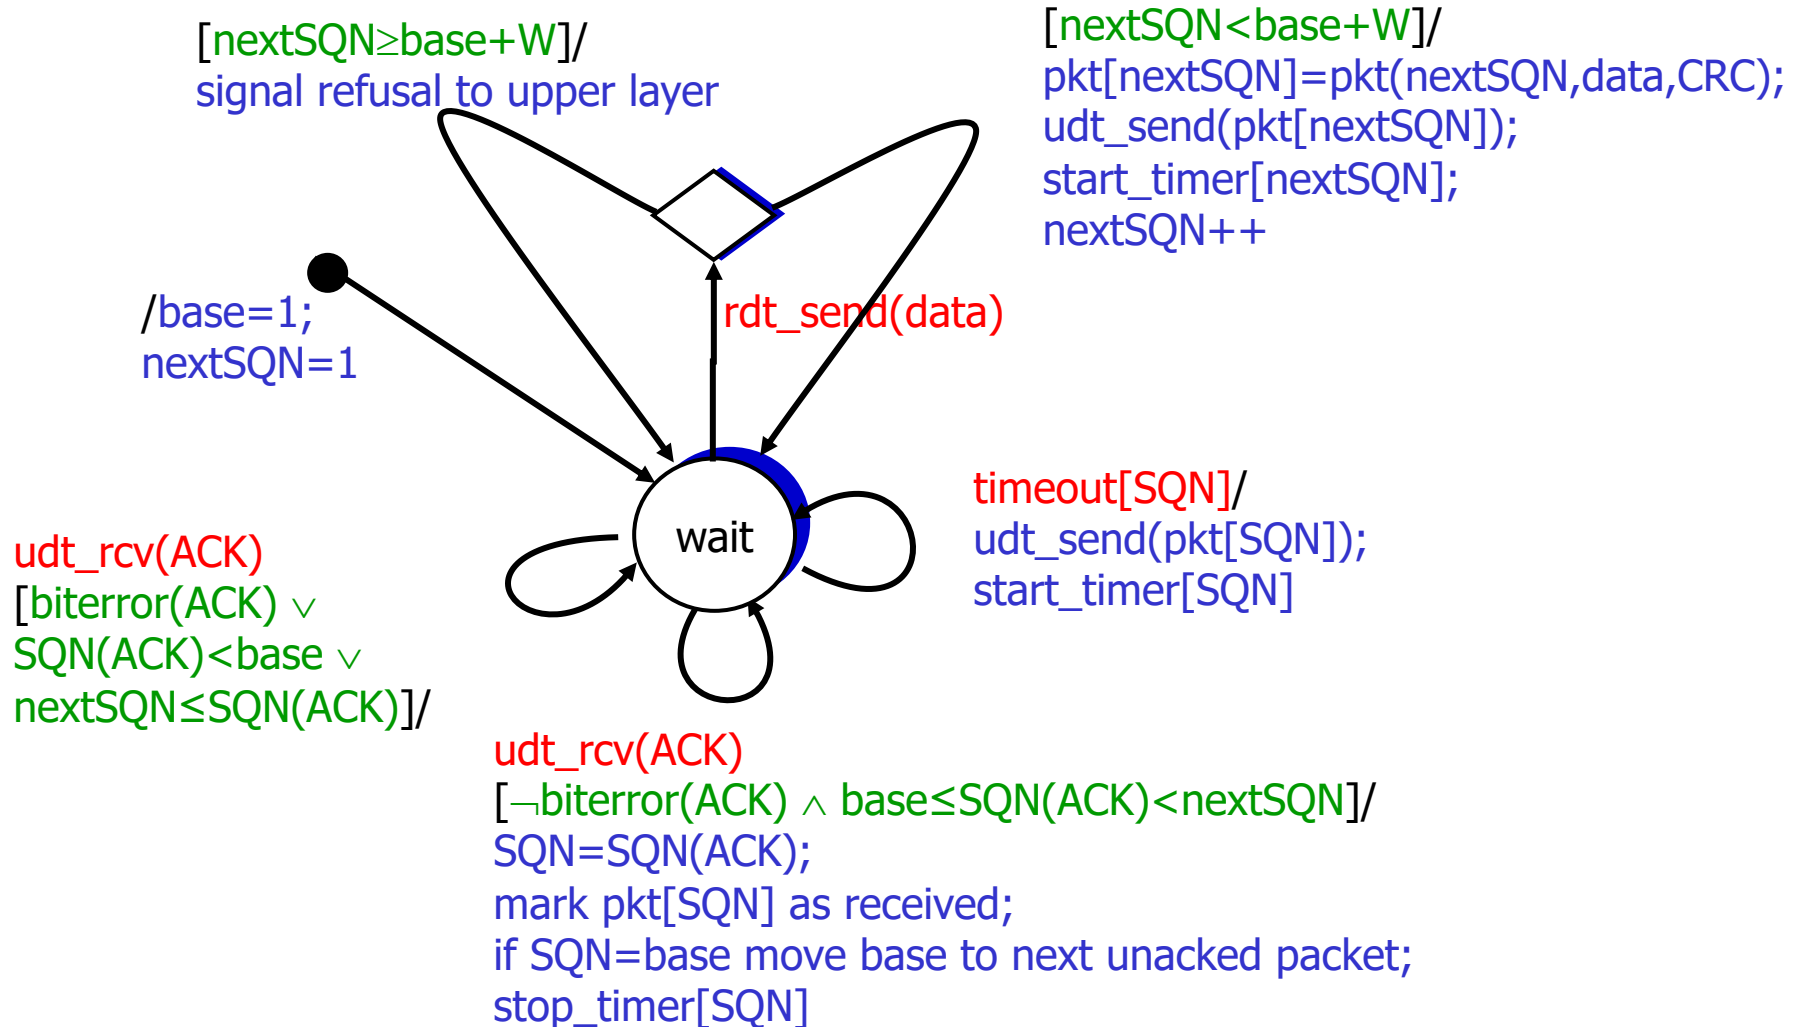
■ Sender

1. If data is available and there is space in the window: send packet, start timer for this packet, increment nextSQN
2. If ACK is received with SQN within window, mark packet with this SQN as acknowledged, stop this timer; if possible, move sliding window
3. If timeout, re-send the associated packet, restart timer


■ Receiver

- If received packet is correct, SQN is within window, send ACK, buffer packet, and, if possible, move sliding window
- If SQN of received packet is from within the previous window, re-send ACK

Selective Repeat: Sender



Selective Repeat: Receiver

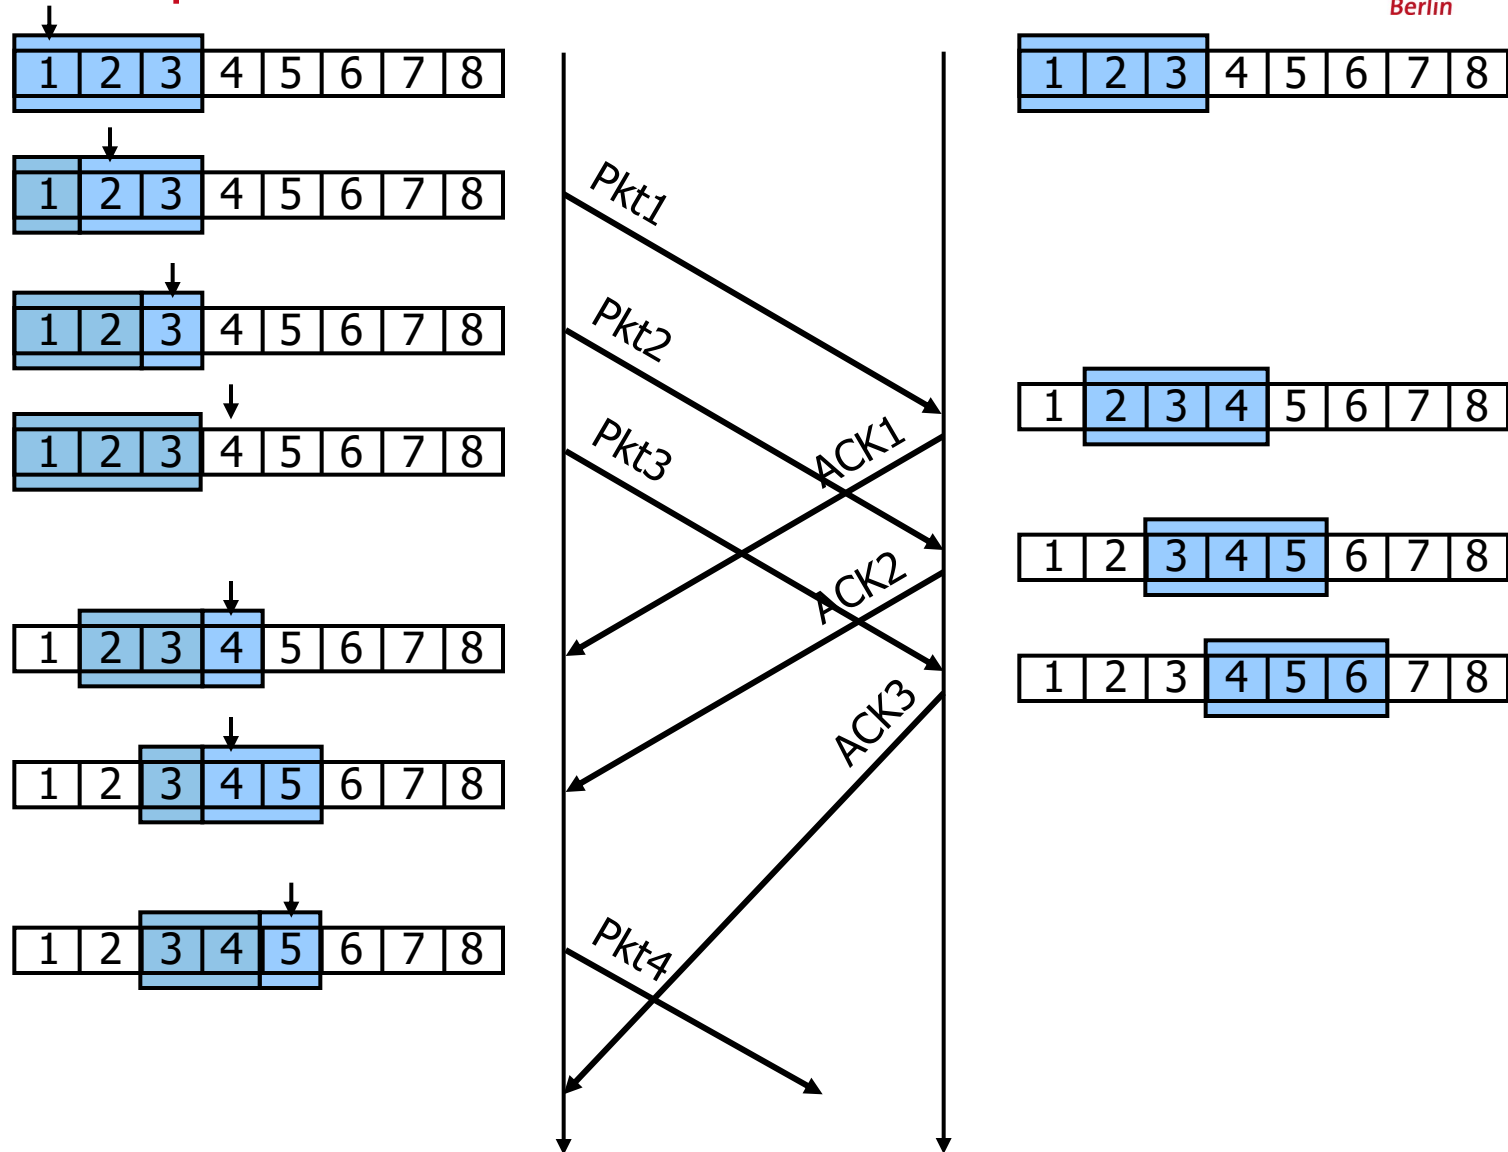


udt_rcv(pkt)
 $[\text{biterror(pkt)} \vee$
 $\text{SQN(pkt) not in [base-W, base+W-1]}]$

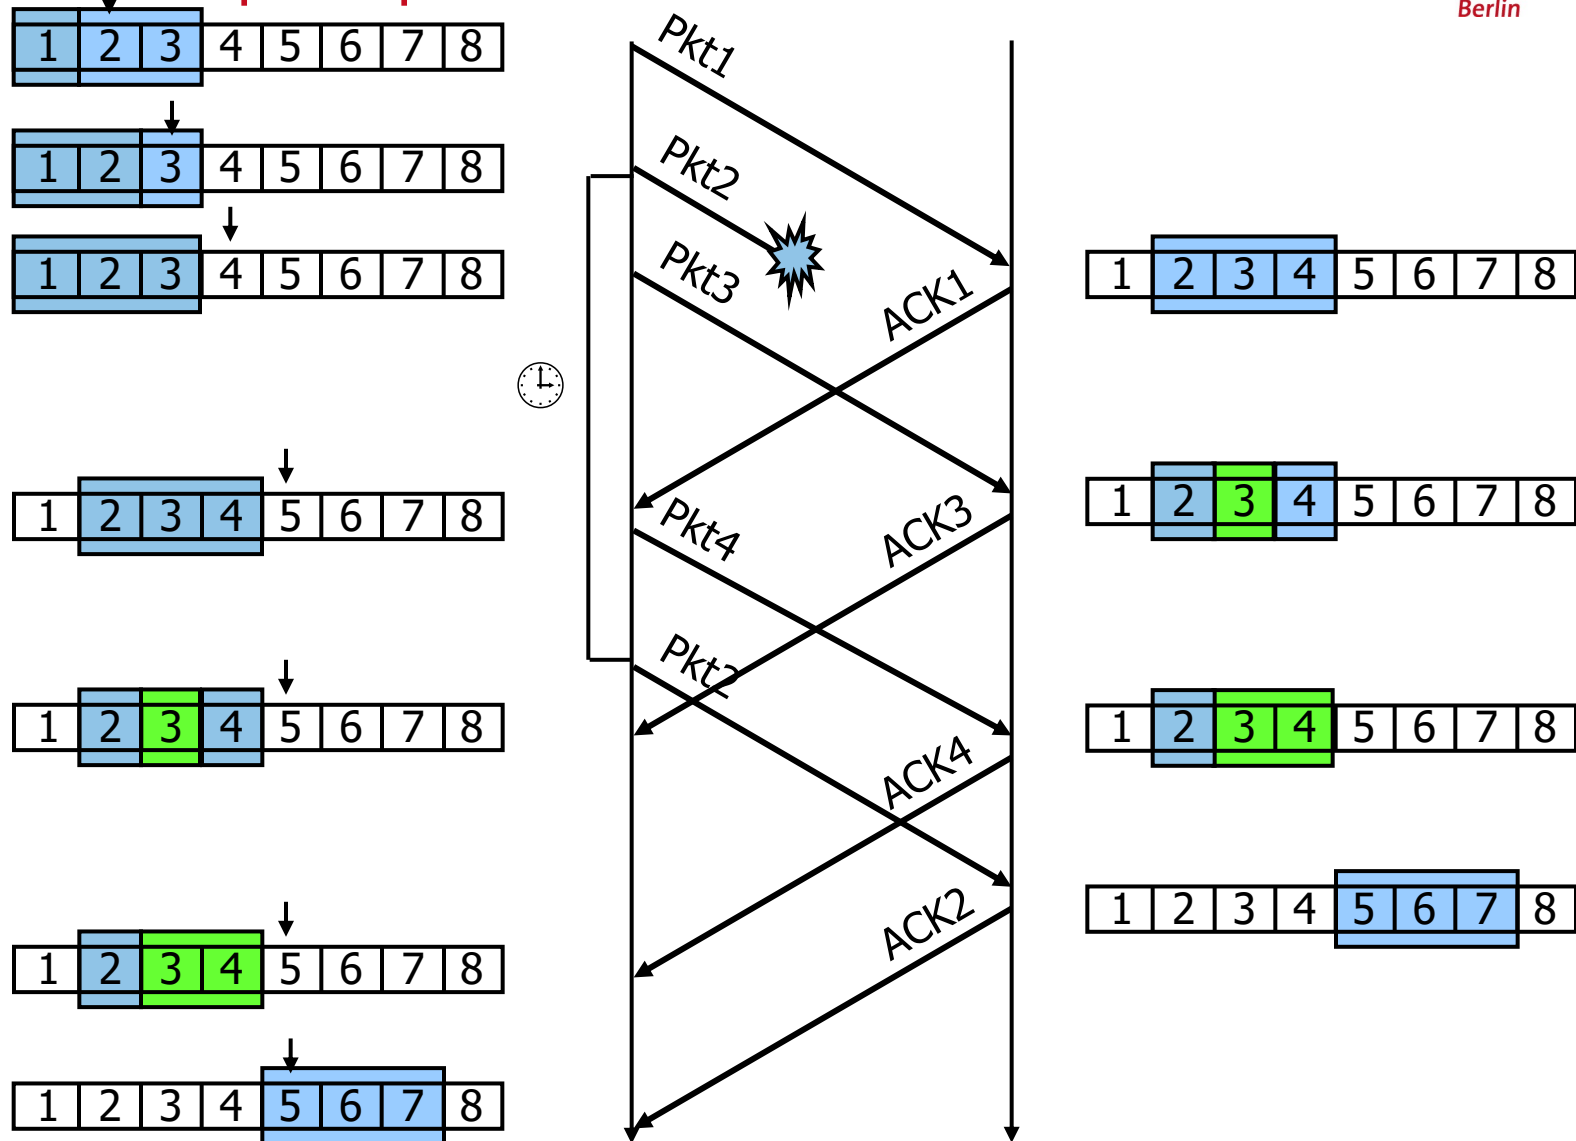
udt_rcv(pkt)
 $[\neg \text{biterror(pkt)} \wedge \text{SQN(pkt) in [base, base+W-1]}]$
 $\text{data} = \text{extractdata(pkt)}$
 buffer data
 $\text{if SQN(pkt) = base}$
 $\quad \text{base} = \text{lowest unrcvd sequence number}$
 $\quad \text{deliver all data until base-1}$
 $\text{ACK} = \text{ACK(SQN(pkt), CRC)}$
 udt_send(ACK)

udt_rcv(pkt)
 $[\neg \text{biterror(pkt)} \wedge \text{SQN(pkt) in [base-W, base-1]}]$
 $\text{ACK} = \text{ACK(SQN(pkt), CRC)}$
 udt_send(ACK)

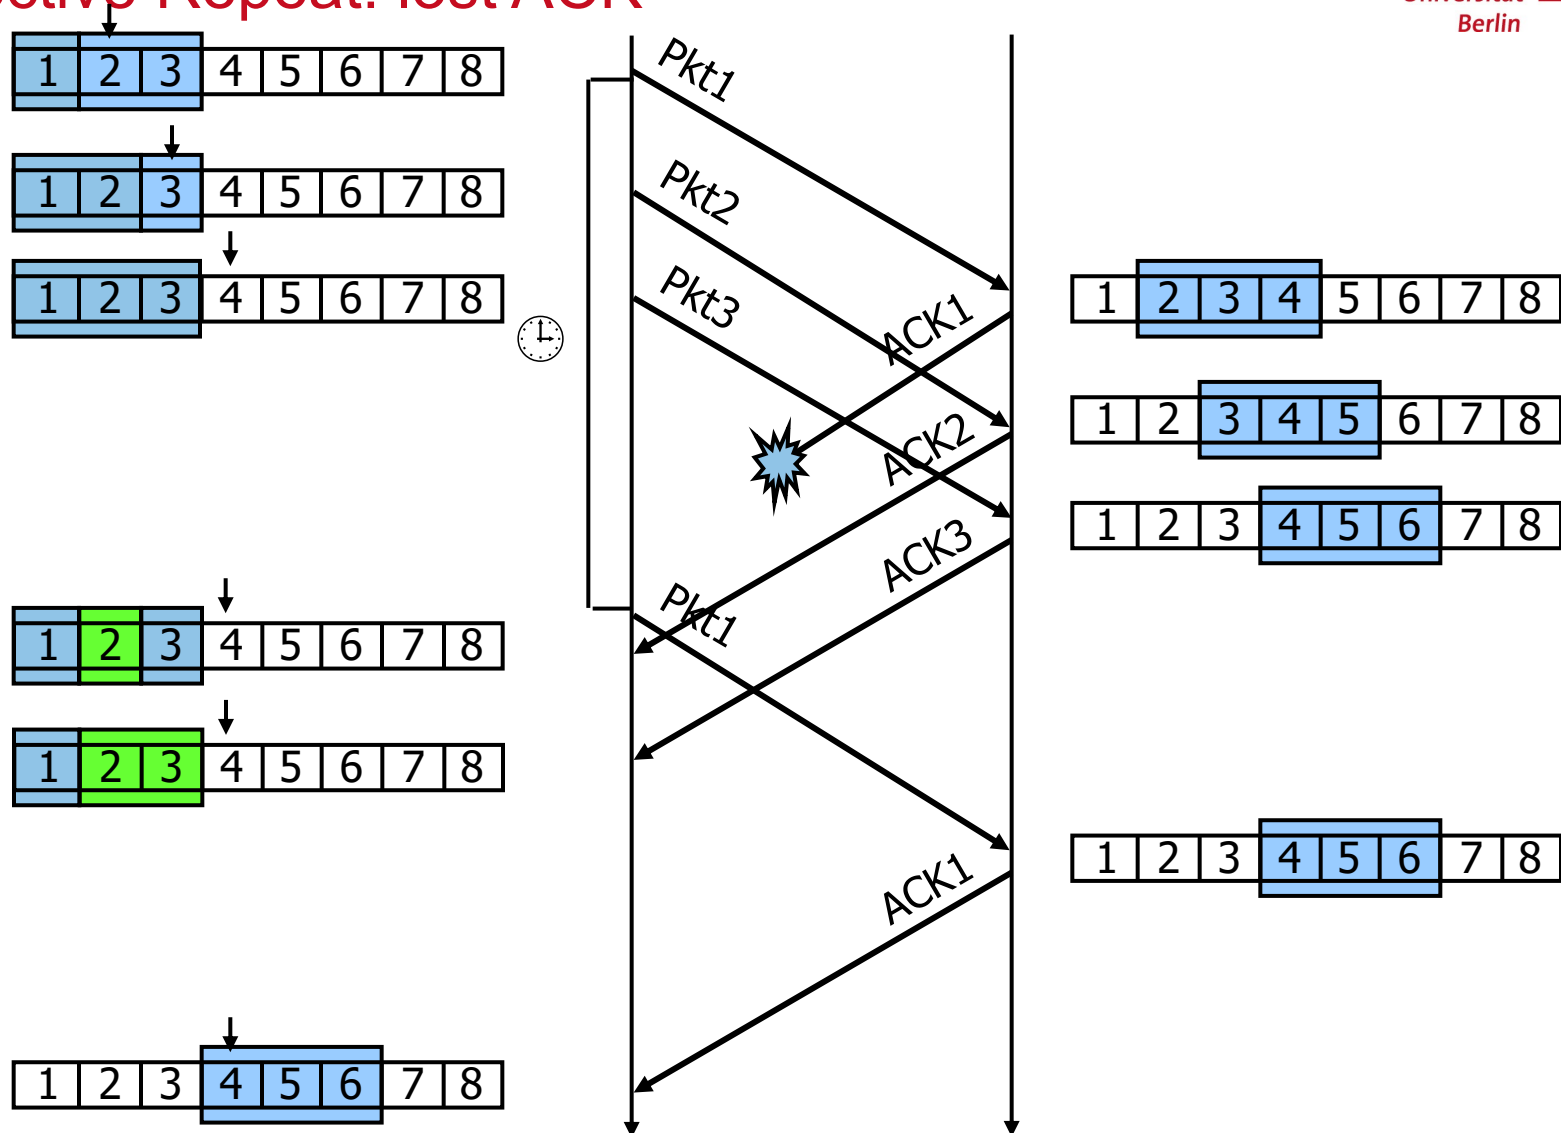
Selective Repeat: normal execution



Selective Repeat: packet loss



Selective Repeat: lost ACK



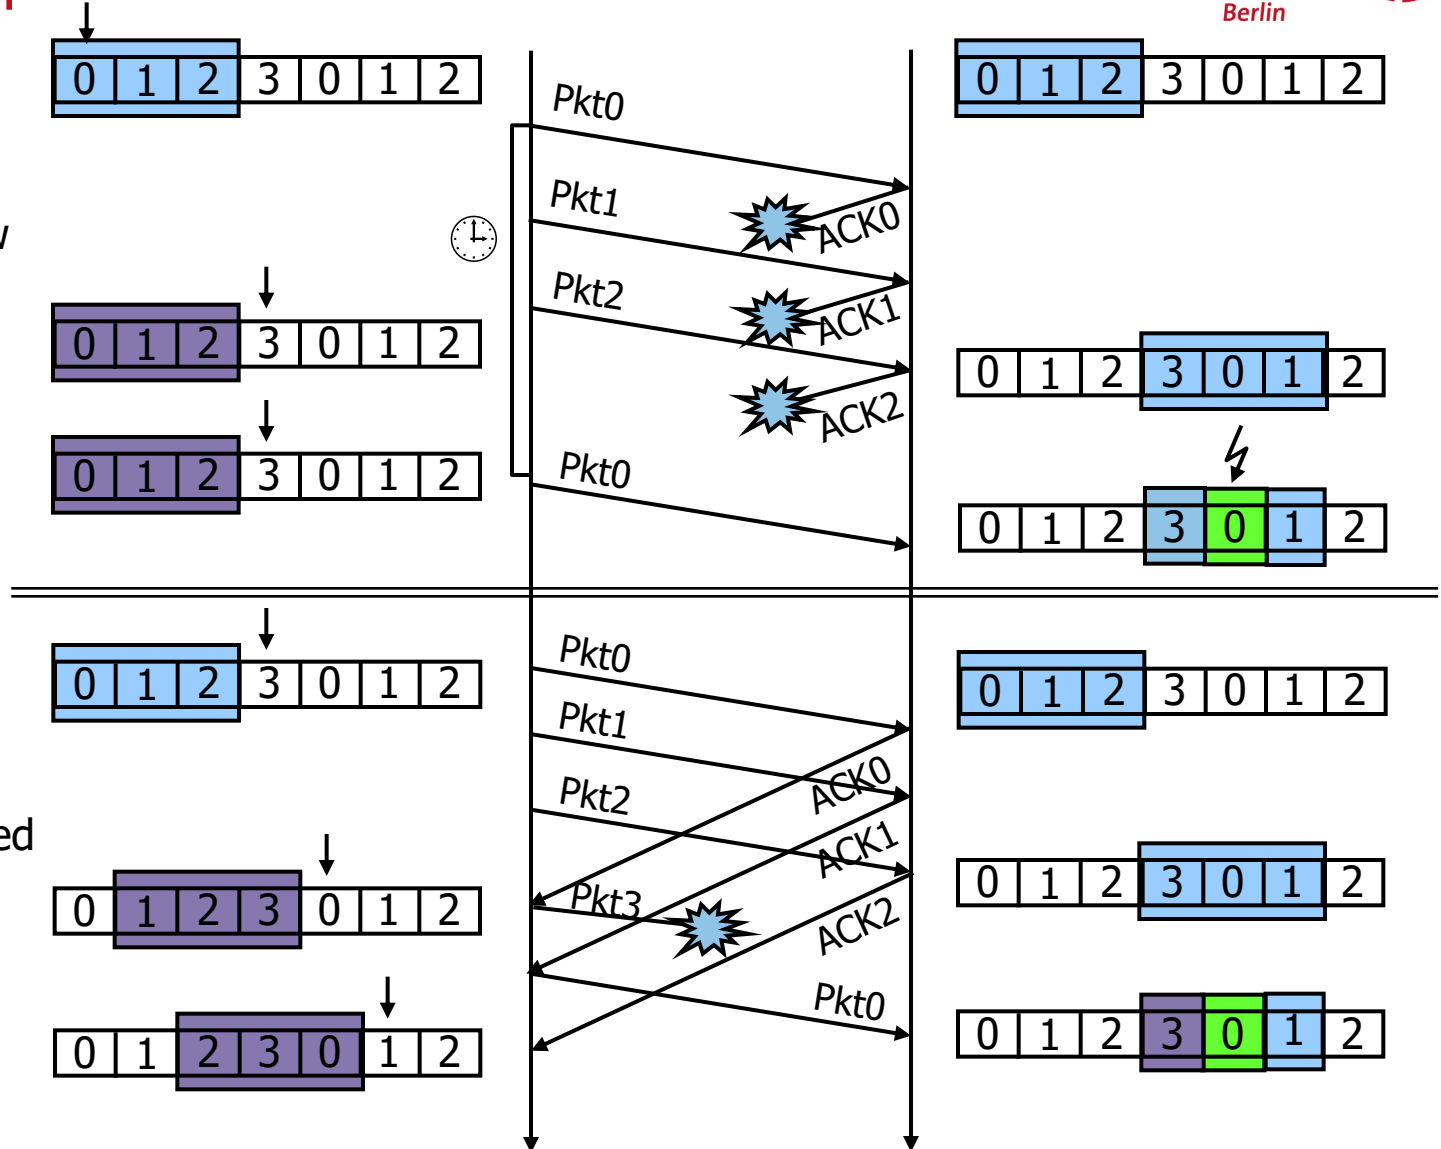
Sequence numbers for sliding window protocols

- Finite sequence numbers with m values
- Cyclic re-use of SQN
- Different packets with same SQN need to be distinguished
- Sufficient condition:
 - If receive window = 1: $W < m$
 - If send window = receive window = $W > 1$:
 $W < (m+1)/2$
- Example for too small sequence number space
 - Sequence numbers $m = 4$, window size $W = 3$
 - $W > (m+1)/2$
 - Receiver cannot distinguish between old and new packets (see next slide)

Selective Repeat

Two possible sequences for $m=4$ and $W=3$,
which the receiver cannot distinguish
and which lead to different results

Re-sent packet 0
is accepted as new
packet



New packet 0
is correctly accepted
as new

Comparison Go-Back-N and Selective Repeat

■ Advantages Go-Back-N

- Cumulative ACKs compensate for lost and delayed ACKs, without the need to re-send packets
- Sender only requires one timer
- Receiver does not need buffer
- Sender and receiver are easy to implement as no vacant spots in the window need to be maintained

■ Advantages Selective Repeat

- Less repeated packets because only really lost packets need to be retransmitted

Performance analysis

Performance analysis

- Possible questions
 - When does Stop-and-Wait block, what's the impact on throughput?
 - How large does the sliding window need to be to optimize throughput?
 - Which is more efficient: Go-Back-N or Selective Repeat?

- Following examples
 - W. Stallings: *Computer Networking with Internet Protocols and Technology*, Pearson Education, 2004 und W. Stallings: *High-Speed Networks, TCP/IP and ATM Design Principles*, Prentice Hall, 1998
 - Typical performance analysis in communication systems
 - Many assumptions to ease analysis
 - Mathematical terms are not that bad 😊
 - More detailed evaluation possible in computer simulation

Bandwidth-delay-product


- A gentle reminder (see Chapter 1)
- **Storage capacity of a channel** (in packets, for packet size L)
 - $a = \frac{RD}{L} = \frac{d/v}{L/R} = \frac{\text{Propagation delay}}{\text{Packet transmission time}}$

(Number of packets sent while the first bit travels from sender to receiver)

Performance analysis: Stop-and-Wait

■ Stop-and-Wait without errors

- (we ignore the time to send ACKs as well as all processing time – this assumption does not significantly change results and eases calculation)



- Throughput in bit per second

$$\text{Throughput} = \frac{L}{\frac{L}{R} + 2D}$$

- Normalized by bitrate (for better comparability)

$$\text{Normalized throughput} = S = \frac{L}{\frac{L}{R} + 2D} \frac{1}{R} = \frac{1}{1 + 2RD/L} = \frac{1}{1 + 2a}$$

- $S = \frac{1}{1+2a} \Rightarrow$ low throughput for large a (channel cannot be filled with packets)

Performance analysis: Stop-and-Wait

- Stop-and-Wait with errors
 - Retransmission after error (timeout)
 - Assumption (again): errors appear with probability p , errors are independent
 - Timeout = $2D$
 - N is the mean number of transmissions of each packet
- Throughout in bit per second

$$\text{Throughout} = \frac{L}{N \left(\frac{L}{R} + 2D \right)}$$

- Normalized by bitrate (for better comparability)

$$\text{Normalized throughput} = S = \frac{L/R}{N \left(\frac{L}{R} + 2D \right)} = \frac{1}{N(1 + 2RD/L)} = \frac{1}{N(1 + 2a)}$$

Performance analysis: Stop-and-Wait

- Calculation of N:
 - Probability that a packet needs to be sent i times is equal to the probability that a transmission failed $i-1$ times followed by a successful attempt
 - $\Pr[i \text{ transmission attempts}] = p^{i-1} \cdot (1-p)$
 - This is a **geometric distribution**, expected value:

$$N = E[\text{transmission attempts}] = \sum_{i=1}^{\infty} i \cdot \Pr[i \text{ transmission attempts}]$$

$$N = \sum_{i=1}^{\infty} i \cdot p^{i-1} \cdot (1-p) = (1-p) \sum_{i=1}^{\infty} i \cdot p^{i-1} = (1-p) \frac{1}{(1-p)^2} = \frac{1}{1-p}$$


$$\text{using } \sum_{i=1}^{\infty} iX^{i-1} = \frac{1}{(1-X)^2} \text{ for } (-1 < X < 1)$$

- Normalized throughput:

$$S = \frac{1-p}{1+2a} \Rightarrow \text{low throughput for large } a \text{ and large } p$$

Performance analysis: Stop-and-Wait

- Normalized throughput of Stop-and-Wait as a function of a:
Stop-and-Wait



For large a and for large p , the normalized throughput gets smaller

Performance analysis: sliding window protocols

- Sliding window protocols without error

- Window size W and packet length L
- Case 1: window is large enough to continuously send before first ACK:

$$W \geq \frac{L/R + 2D}{L/R} = 1 + 2a$$

$$S = \frac{W \cdot L}{W \cdot L/R} \cdot \frac{1}{R} = 1$$

- Case 2: window is too small:


$$W < 1 + 2a$$

$$S = \frac{W \cdot L}{L/R + 2D} \cdot \frac{1}{R} = \frac{W}{1 + 2a}$$

$$S = \begin{cases} 1 & W \geq 1 + 2a \\ \frac{W}{1 + 2a} & W < 1 + 2a \end{cases}$$

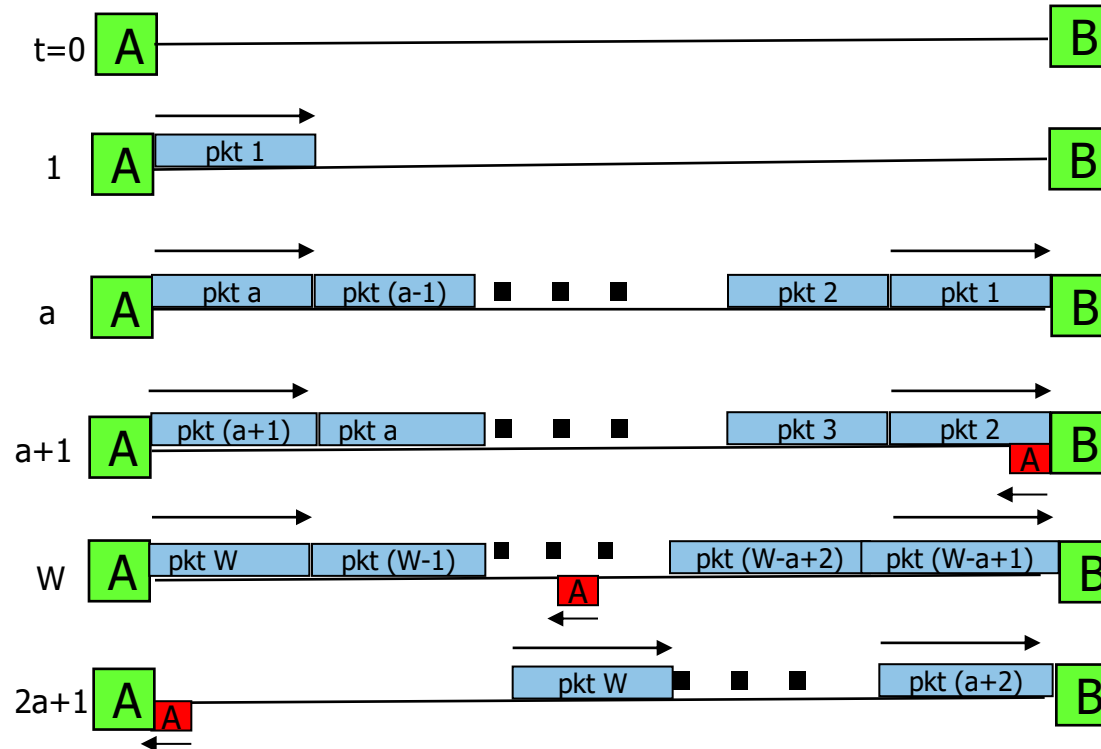
Performance analysis: sliding window protocols

■ Schematic procedure



Performance analysis: sliding window protocols


■ Schematic procedure



$$W < 2a + 1$$

Performance analysis: sliding window protocols

- Normalized throughput of sliding window protocols as a function of a :



→ Window size W
Is sufficient to
fill the pipe, if
 $a = (W-1)/2$

Performance analysis: sliding window protocols

■ Selective Repeat with errors

- Assumption: errors with probability p are independent
- $N = E[\text{transmission attempts}] = 1/(1-p)$
- Throughput in error-free case needs to be divided by N :

$$S = \begin{cases} \frac{1}{N} = \frac{1}{1/(1-p)} = 1-p & W \geq 1 + 2a \\ \frac{W}{N \cdot (1 + 2a)} = \frac{W}{1/(1-p) \cdot (1 + 2a)} = \frac{W(1-p)}{1 + 2a} & W < 1 + 2a \end{cases}$$

$$S = \begin{cases} 1-p & W \geq 1 + 2a \\ \frac{W(1-p)}{1 + 2a} & W < 1 + 2a \end{cases}$$

Performance analysis: sliding window protocols

■ Go-back-N with errors

- Every error requires retransmission of K packets
- Assumption: in case of error, the pipe is filled and all packets in the window needs to be retransmitted:

$$K = \begin{cases} 1 + 2a & W \geq 1 + 2a \\ W & W < 1 + 2a \end{cases}$$

- If the erroneous packet needs to be retransmitted i times, in total $1+(i-1)K = (1-K)+Ki$ packets need to be sent

$$\begin{aligned} N &= \sum_{i=1}^{\infty} ((1-K) + Ki) \cdot p^{i-1} \cdot (1-p) = (1-K)(1-p) \sum_{i=1}^{\infty} p^{i-1} + K(1-p) \sum_{i=1}^{\infty} i \cdot p^{i-1} \\ &= (1-K)(1-p) \frac{1}{(1-p)} + K(1-p) \frac{1}{(1-p)^2} \\ &= 1-K + \frac{K}{1-p} = \frac{1-p+Kp}{1-p} \\ &\text{using } \sum_{i=1}^{\infty} X^{i-1} = \frac{1}{1-X} \text{ for } (-1 < X < 1) \end{aligned}$$

Performance analysis: sliding window protocols

■ Given K, we get:

$$N = \begin{cases} \frac{1-p+Kp}{1-p} = \frac{1-p+(1+2a)p}{1-p} = \frac{1+2ap}{1-p} & W \geq 1+2a \\ \frac{1-p+Kp}{1-p} = \frac{1-p+Wp}{1-p} & W < 1+2a \end{cases}$$


- Normalized throughput (dividing throughput w/o errors by N):

$$S = \begin{cases} \frac{1}{N} = \frac{1-p}{1+2ap} & W \geq 1+2a \\ \frac{W}{N \cdot (1+2a)} = \frac{W(1-p)}{(1-p+Wp) \cdot (1+2a)} & W < 1+2a \end{cases}$$

$$S = \begin{cases} \frac{1-p}{1+2ap} & W \geq 1+2a \\ \frac{W(1-p)}{(1-p+Wp) \cdot (1+2a)} & W < 1+2a \end{cases}$$

Performance analysis: sliding window protocols


- Normalized throughput of G-Back-N and Selective Repeat as a function of a , $p = 10^{-3}$:



→ only for large window sizes, we notice an advantage of Selective Repeat

Performance analysis: sliding window protocols


- Normalized throughput of G-Back-N and Selective Repeat as a function of a , $p = 10^{-2}$:



→ due to the higher loss rate, we notice the advantage of Selective Repeat also for smaller window sizes

Performance analysis: sliding window protocols


- Normalized throughput of G-Back-N and Selective Repeat as a function of a , $p = 10^{-1}$:



→ for high loss rates, the advantage of Selective Repeat is very clearly visible

Performance analysis: sliding window protocols


- Normalized throughput of G-Back-N and Selective Repeat as a function of W , $p = 10^{-3}$:



→ for large a and for large window sizes, there is a significant difference between Go-Back-N and Selective Repeat

Performance analysis: sliding window protocols


- Normalized throughput of G-Back-N and Selective Repeat as a function of W , $p = 10^{-2}$:



→ the difference grows for higher loss rates

Performance analysis: sliding window protocols

- Normalized throughput of G-Back-N and Selective Repeat as a function of W , $p = 10^{-1}$:



→ and even more