

Cognitive Algorithms - Exercise 5

Unsupervised Learning

Daniel Wujecki

December 20, 2020

Organizational Remarks

- Grading of exercise 2 is not finished
 - Many submissions
 - Hopefully done within the next days
- Teaching evaluation in next week, also for elective courses
- please send mails **ONLY** from your **TU-Mail**

Today's Tutorial

- Due to your feedback no webcam - focus on the content ;)
- Repetition of theoretical concepts
 - Proofs are shown in the lecture videos

Today's Tutorial

- Due to your feedback no webcam - focus on the content ;)
- Repetition of theoretical concepts
 - Proofs are shown in the lecture videos
- If too many repetitions for you, feel free to skip!
 - Since you can skip, I am quite detailed for those who have more difficulties
- Many practical examples to visualize the concepts
 - Code will not be provided (contains Quiz solutions)
 - You do not need to understand the code I show

- 1 Task 1 + 4 - Eigenvalues and Eigenvectors
- 2 Task 2 + 3 + 4 - Principal Component Analysis
- 3 Task 5 - Non-negative matrix factorization
- 4 K-Means Clustering

- 1 Task 1 + 4 - Eigenvalues and Eigenvectors
- 2 Task 2 + 3 + 4 - Principal Component Analysis
- 3 Task 5 - Non-negative matrix factorization
- 4 K-Means Clustering

Repetition: Eigenvalues and Eigenvectors

Given a Matrix $A \in \mathbb{R}^{d \times d}$, then a non-zero vector $\mathbf{v} \in \mathbb{C}^d \setminus \{\mathbf{0}\}$ is called an eigenvector, if there is an eigenvalue $\lambda \in \mathbb{C}$, such that

$$A\mathbf{v} = \lambda\mathbf{v}$$

Repetition: Eigenvalues and Eigenvectors

Given a Matrix $A \in \mathbb{R}^{d \times d}$, then a non-zero vector $\mathbf{v} \in \mathbb{C}^d \setminus \{\mathbf{0}\}$ is called an eigenvector, if there is an eigenvalue $\lambda \in \mathbb{C}$, such that

$$A\mathbf{v} = \lambda\mathbf{v}$$

- eigenvectors are special directions, for which a matrix A is only scales, but not rotate
- A matrix is singular if one or more eigenvalues are zero

$$\lambda = 0 \Leftrightarrow \text{Kern}(A) \neq \emptyset$$

Repetition: Scaled Eigenvectors

If $\mathbf{v} \in \mathbb{R}^d$ is a eigenvector of $A \in \mathbb{R}^{d \times d}$, then also $\alpha \mathbf{v}$, $\alpha \in \mathbb{R}$, $\alpha \neq 0$ is a eigenvector (Task 4.1).

Proof

$$A(\alpha \mathbf{v}) = \alpha(A\mathbf{v}) = \alpha(\lambda \mathbf{v}) = \lambda(\alpha \mathbf{v})$$

where we used the homogeneity of linear mappings.

Repetition: Eigenvalues for special matrices

upper and lower triangular matrices

$$\begin{pmatrix} \lambda_1 & a_{1,2} & \cdots & a_{1,n} \\ & \lambda_2 & \cdots & a_{2,n} \\ & & \ddots & \vdots \\ 0 & & & \lambda_n \end{pmatrix}$$

diagonal matrices

$$\begin{pmatrix} \lambda_1 & & & 0 \\ & \lambda_2 & & \\ & & \ddots & \\ 0 & & & \lambda_n \end{pmatrix}$$

\Rightarrow eigenvalues are on the diagonal (Task 1.2, 1.3)

Rotation Matrices I

What about rotation matrices?

$$R_\theta = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix}$$

Rotation Matrices I

What about rotation matrices?

$$R_\theta = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix}$$

- They rotate vectors by an angle θ , but do not scale any vector
⇒ no real eigenvalues, but complex eigenvalues (Task 1.4)

Rotation Matrices I

What about rotation matrices?

$$R_\theta = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix}$$

- They rotate vectors by an angle θ , but do not scale any vector
 \Rightarrow no real eigenvalues, but complex eigenvalues (Task 1.4)
- If $\theta = k\pi$, $k \in \mathbb{Z}$ they scale by $\lambda = \pm 1$
 - special case with real eigenvalues (Task 1.5)

Rotation Matrices II

- Rotation matrices R_θ are always orthogonal with $\det |R_\theta| = +1$
- Are all orthogonal matrices also rotation matrices? (Task 4.2)

Rotation Matrices II

- Rotation matrices R_θ are always orthogonal with $\det |R_\theta| = +1$
- Are all orthogonal matrices also rotation matrices? (Task 4.2)
- For a orthogonal matrix V the determinant is $\det |V| = \pm 1$

$$V = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

- V is orthogonal, since $V^\top V = I$, but $\det |V| = -1$ and thus it is not a rotation matrix

Repetition: Eigendecomposition

- If a matrix $A \in \mathbb{R}^{d \times d}$ is symmetric
 - all eigenvalues $\lambda_1, \dots, \lambda_d$ are real
 - all eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_d$ are orthogonal to each other (Task 1.1)

Repetition: Eigendecomposition

- If a matrix $A \in \mathbb{R}^{d \times d}$ is symmetric
 - all eigenvalues $\lambda_1, \dots, \lambda_d$ are real
 - all eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_d$ are orthogonal to each other (Task 1.1)
- For symmetric matrices there is a decomposition $A = U\Lambda U^T$ into
 - a diagonal matrix Λ with the eigenvalues on the diagonal
 - a orthogonal matrix U with the eigenvectors in the columns

- 1 Task 1 + 4 - Eigenvalues and Eigenvectors
- 2 Task 2 + 3 + 4 - Principal Component Analysis
- 3 Task 5 - Non-negative matrix factorization
- 4 K-Means Clustering

Dimensionality of Data

- Given some high dimensional data
 $X = [\mathbf{x}_1, \dots, \mathbf{x}_n] \in \mathbb{R}^{d \times n}$
- Naturally, not all dimensions contain the same amount of information
- Extrem Example: if $d \gg n$, samples can not span d -dimensional vector space, even if all samples are linearly independent

Dimensionality of Data

- Given some high dimensional data
 $X = [\mathbf{x}_1, \dots, \mathbf{x}_n] \in \mathbb{R}^{d \times n}$
- Naturally, not all dimensions contain the same amount of information
- Extrem Example: if $d \gg n$, samples can not span d-dimensional vector space, even if all samples are linearly independent



Dimensionality of Data

- Given some high dimensional data
 $X = [\mathbf{x}_1, \dots, \mathbf{x}_n] \in \mathbb{R}^{d \times n}$
- Naturally, not all dimensions contain the same amount of information
- Extrem Example: if $d \gg n$, samples can not span d-dimensional vector space, even if all samples are linearly independent
- How do we measure the 'amount of information'?



Principle Component Analysis



- Which dimensions contain the most information?


⇒ directions \mathbf{w} with the highest variance

$$\operatorname{argmax}_{\mathbf{w}} \operatorname{var}(\mathbf{w}^{\top} X) = \operatorname{argmax}_{\mathbf{w}} \mathbf{w}^{\top} \Sigma_X \mathbf{w}$$

⇒ directions \mathbf{w} that minimizes the noise

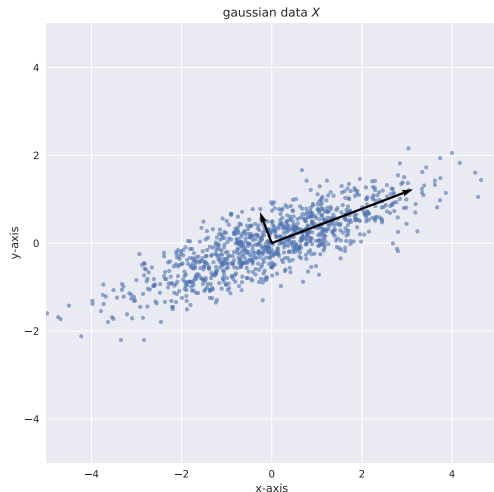
$$\operatorname{argmin}_{\mathbf{w}} \sum_{i=1}^n \|(\mathbf{w}^{\top} \mathbf{x}_i) \mathbf{w} - \mathbf{x}_i\|^2 = \operatorname{argmax}_{\mathbf{w}} \mathbf{w}^{\top} \Sigma_X \mathbf{w}$$

Principle Component Analysis




- $\mathbf{w}^T \Sigma_X \mathbf{w}$ is maximized by the eigenvector of Σ_X with the largest eigenvalue
- The eigenvectors of Σ_X are the principle components

Principle Component Analysis



- $\mathbf{w}^T \Sigma_X \mathbf{w}$ is maximized by the eigenvector of Σ_X with the largest eigenvalue
- The eigenvectors of Σ_X are the principle components
- PCs can be constrained to unit length $\|\mathbf{w}\| = 1$

Principle Component Analysis



- $\mathbf{w}^T \Sigma_X \mathbf{w}$ is maximized by the eigenvector of Σ_X with the largest eigenvalue
- The eigenvectors of Σ_X are the principle components
- PCs can be constrained to unit length $\|\mathbf{w}\| = 1$
- Variance along PC is given by the corresponding eigenvalue, since

$$\text{var}(\mathbf{w}^T X) = \mathbf{w}^T \Sigma_X \mathbf{w} = \mathbf{w}^T \lambda \mathbf{w} = \lambda \mathbf{w}^T \mathbf{w} = \lambda$$

Dimensionality Reduction with PCA - 3D example

- Given a 3D dataset $X \in \mathbb{R}^{3 \times n}$
- Calculate the EVD of the covariance $\Sigma_X = W\Lambda W^T$
- What happens when we map the data onto the PC?

Dimensionality Reduction with PCA - 3D example

- Given a 3D dataset $X \in \mathbb{R}^{3 \times n}$
- Calculate the EVD of the covariance $\Sigma_X = W \Lambda W^T$
- What happens when we map the data onto the PC?

$$W^T \mathbf{x} = \begin{bmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \\ \mathbf{w}_3^T \end{bmatrix} \mathbf{x} = \begin{bmatrix} \mathbf{w}_1^T \mathbf{x} \\ \mathbf{w}_2^T \mathbf{x} \\ \mathbf{w}_k^T \mathbf{x} \end{bmatrix} = \mathbf{x}' = \begin{bmatrix} x'_x \\ x'_y \\ x'_z \end{bmatrix}$$

- Data is decorrelated: PCs become aligned with the coordinate axes
 - PCs are orthogonal to each other
 - orthogonal matrix W corresponds to a rotation

Dimensionality Reduction with PCA - 3D example

- We could drop now the dimension with lowest variance
- For the decorrelated data that is just dropping one coordinate axis

$$\begin{bmatrix} x'_x \\ x'_y \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^\top \mathbf{x} \\ \mathbf{w}_2^\top \mathbf{x} \end{bmatrix} \in \mathbb{R}^2$$

Dimensionality Reduction with PCA - 3D example

- We could drop now the dimension with lowest variance
- For the decorrelated data that is just dropping one coordinate axis

$$\begin{bmatrix} x'_x \\ x'_y \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^\top \mathbf{x} \\ \mathbf{w}_2^\top \mathbf{x} \end{bmatrix} \in \mathbb{R}^2$$

- To restore the same information we have to map the data back

$$[\mathbf{w}_1, \mathbf{w}_2] \begin{pmatrix} x'_x \\ x'_y \end{pmatrix} = x'_x \cdot \mathbf{w}_1 + x'_y \cdot \mathbf{w}_2 \approx \begin{bmatrix} x_x \\ x_y \\ x_z \end{bmatrix}$$

PCA Algorithm

Algorithm 1: Principal Component Analysis

Input: Dataset $X \in \mathbb{R}^{d \times n}$; Number of PCs k

Output: First k PCs $W \in \mathbb{R}^{d \times k}$, Hidden causes $H \in \mathbb{R}^{k \times n}$ of the dataset

- 1 Compute covariance $\Sigma_X = \frac{1}{n}(X - \bar{X})(X - \bar{X})^\top$
 - 2 Compute EVD $\Sigma_X = V\Lambda V^\top$
 - 3 Take first k eigenvectors corresponding to largest eigenvalues $W = [\mathbf{v}_1, \dots, \mathbf{v}_k]$
 - 4 Project data $H = W^\top X$
 - 5 **return** W and H
-

Task 2.1 - PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace

Task 2.1 - PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace
 - $\bar{X} = \mathbf{0}$

Task 2.1 - PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace

① $\bar{X} = \mathbf{0}$

② $\Sigma_X = \frac{1}{2} \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$

Task 2.1 - PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace

$$\textcircled{1} \bar{X} = \mathbf{0}$$

$$\textcircled{2} \Sigma_X = \frac{1}{2} \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$\textcircled{3} \mathbf{w}^T X = [1, 0, 0] \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} = [-1, 1]$$

Task 2.1 - PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$

- Project the data onto a 1-dimensional subspace

① $\bar{X} = \mathbf{0}$

② $\Sigma_X = \frac{1}{2} \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$

③ $\mathbf{w}^T X = [1, 0, 0] \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} = [-1, 1]$

- Linearly dependent samples \Rightarrow Representation makes sense

Kernel PCA

- If dimensionality becomes larger than number of samples $d \gg n$
 - Only $\leq n$ non-zero eigenvalues
 - Covariance Σ_X becomes singular, $\text{Rank}(\Sigma_X) \leq n$
 - Covariance is positive *semi* definite
 - Covariance $\Sigma_X \in \mathbb{R}^{d \times d}$ will be very large
 - High time and space complexity for calculation
- Kernel PCA should be used to estimate n first Principal Components

Kernel PCA Algorithm

Algorithm 2: Kernel PCA

Input: Dataset $X \in \mathbb{R}^{d \times n}$ with $d \gg n$; Number of PCs k

Output: First k PCs $W \in \mathbb{R}^{d \times k}$, Hidden causes $H \in \mathbb{R}^{k \times n}$ of the dataset

- 1 Compute Kernel $K = (X - \bar{X})^\top (X - \bar{X}) \in \mathbb{R}^{n \times n}$
 - 2 Compute EVD $K = V \Lambda V^\top$
 - 3 Take first k eigenvectors corresponding to largest eigenvalues $\alpha = [\mathbf{v}_1, \dots, \mathbf{v}_k]$
 - 4 Calculate $W = X\alpha \in \mathbb{R}^{d \times k}$
 - 5 Project data $H = W^\top X$
 - 6 **return** W and H
-

Task 2.2 - Kernel PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace using Kernel PCA

Task 2.2 - Kernel PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace using Kernel PCA
 - 1 $\bar{X} = \mathbf{0}$

Task 2.2 - Kernel PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace using Kernel PCA

① $\bar{X} = \mathbf{0}$

② $K = \begin{bmatrix} -1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & -1 \\ -1 & 1 \end{bmatrix} = \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix}$

Task 2.2 - Kernel PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace using Kernel PCA

$$\textcircled{1} \bar{X} = \mathbf{0}$$

$$\textcircled{2} K = \begin{bmatrix} -1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & -1 \\ -1 & 1 \end{bmatrix} = \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix}$$

$$\textcircled{3} \mathbf{w} = X\alpha = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1 \\ 1 \end{bmatrix} = \begin{bmatrix} 2 \\ 0 \\ 0 \end{bmatrix} \Rightarrow \frac{\mathbf{w}}{\|\mathbf{w}\|} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$$

Task 2.2 - Kernel PCA Algorithm

- Consider a dataset with two samples $X = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix}$
- Project the data onto a 1-dimensional subspace using Kernel PCA

$$\textcircled{1} \bar{X} = \mathbf{0}$$

$$\textcircled{2} K = \begin{bmatrix} -1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & -1 \\ -1 & 1 \end{bmatrix} = \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} 2 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1 & 1 \\ 1 & 1 \end{bmatrix}$$

$$\textcircled{3} \mathbf{w} = X\alpha = \begin{bmatrix} -1 & 1 \\ 0 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} -1 \\ 1 \end{bmatrix} = \begin{bmatrix} 2 \\ 0 \\ 0 \end{bmatrix} \Rightarrow \frac{\mathbf{w}}{\|\mathbf{w}\|} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$$



- Same result as before (but with smaller time and space complexity)

Task 4.3

- Consider a dataset $X \in \mathbb{R}^{d \times n}$ and a orthogonal rotation matrix $U \in \mathbb{R}^{d \times d}$
- Proof that the covariance Σ_X and the covariance Σ_{UX} of the rotated data have the same eigenvalues and the eigenvectors are also rotated by U

Task 4.3

- Consider a dataset $X \in \mathbb{R}^{d \times n}$ and a orthogonal rotation matrix $U \in \mathbb{R}^{d \times d}$
- Proof that the covariance Σ_X and the covariance Σ_{UX} of the rotated data have the same eigenvalues and the eigenvectors are also rotated by U



Task 4.3

- Consider a dataset $X \in \mathbb{R}^{d \times n}$ and a orthogonal rotation matrix $U \in \mathbb{R}^{d \times d}$
- Proof that the covariance Σ_X and the covariance Σ_{UX} of the rotated data have the same eigenvalues and the eigenvectors are also rotated by U
- We have

$$\Sigma_{UX} = \frac{1}{N} UX(UX)^{\top} = U \frac{1}{N} XX^{\top} U^{\top} = U \Sigma_X U^{\top}$$

Task 4.3

- Consider a dataset $X \in \mathbb{R}^{d \times n}$ and a orthogonal rotation matrix $U \in \mathbb{R}^{d \times d}$
- Proof that the covariance Σ_X and the covariance Σ_{UX} of the rotated data have the same eigenvalues and the eigenvectors are also rotated by U
- We have

$$\Sigma_{UX} = \frac{1}{N} UX(UX)^T = U \frac{1}{N} XX^T U^T = U \Sigma_X U^T$$

Thus

$$\Sigma_X \mathbf{v} = \lambda \mathbf{v} \Leftrightarrow U \Sigma_X \mathbf{v} = \lambda U \mathbf{v} \Leftrightarrow \underbrace{U \Sigma_X U^T}_{\Sigma_{UX}} \underbrace{U \mathbf{v}}_{:= \mathbf{z}} = \lambda \underbrace{U \mathbf{v}}_{= \mathbf{z}}$$

Task 4.3

- Consider a dataset $X \in \mathbb{R}^{d \times n}$ and a orthogonal rotation matrix $U \in \mathbb{R}^{d \times d}$
- Proof that the covariance Σ_X and the covariance Σ_{UX} of the rotated data have the same eigenvalues and the eigenvectors are also rotated by U
- We have

$$\Sigma_{UX} = \frac{1}{N} UX(UX)^T = U \frac{1}{N} XX^T U^T = U \Sigma_X U^T$$

Thus

$$\Sigma_X \mathbf{v} = \lambda \mathbf{v} \Leftrightarrow U \Sigma_X \mathbf{v} = \lambda U \mathbf{v} \Leftrightarrow \underbrace{U \Sigma_X U^T}_{\Sigma_{UX}} \underbrace{U \mathbf{v}}_{:= \mathbf{z}} = \lambda \underbrace{U \mathbf{v}}_{= \mathbf{z}}$$

Thus, if \mathbf{v} is a eigenvector of Σ_X with corresponding eigenvalue λ , $U \mathbf{v}$ is eigenvector of Σ_{UX} with eigenvalue λ .

Task 3

Given two eigenvalues of a covariance matrix, sketch two corresponding 2-dimensional datasets with uncorrelated and correlated features, respectively.

- Reminder: Variance along PC is given by the corresponding eigenvalue, since

$$\text{var}(\mathbf{w}^\top X) = \mathbf{w}^\top \Sigma_X \mathbf{w} = \mathbf{w}^\top \lambda \mathbf{w} = \lambda \mathbf{w}^\top \mathbf{w} = \lambda$$

Task 3.1 - Variances along PCs

Given two eigenvalues of a covariance matrix, sketch two corresponding 2-dimensional datasets with uncorrelated and correlated features, respectively.

$$\lambda_1 = 1$$



$$\lambda_2 = 1$$

Task 3.1 - Variances along PCs

Given two eigenvalues of a covariance matrix, sketch two corresponding 2-dimensional datasets with uncorrelated and correlated features, respectively.

$$\lambda_1 = 1$$

$$\lambda_2 = 1$$



Task 3.2 - Variances along PCs

Given two eigenvalues of a covariance matrix, sketch two corresponding 2-dimensional datasets with uncorrelated and correlated features, respectively.

$$\lambda_1 = 1$$



$$\lambda_2 = 5$$

Task 3.2 - Variances along PCs

Given two eigenvalues of a covariance matrix, sketch two corresponding 2-dimensional datasets with uncorrelated and correlated features, respectively.

$$\lambda_1 = 1$$

$$\lambda_2 = 5$$



Task 3.3 - Variances along PCs

Given two eigenvalues of a covariance matrix, sketch two corresponding 2-dimensional datasets with uncorrelated and correlated features, respectively.

$$\lambda_1 = 1$$



$$\lambda_2 = 0$$

Task 3.3 - Variances along PCs

Given two eigenvalues of a covariance matrix, sketch two corresponding 2-dimensional datasets with uncorrelated and correlated features, respectively.

$$\lambda_1 = 1$$

$$\lambda_2 = 0$$



- 1 Task 1 + 4 - Eigenvalues and Eigenvectors
- 2 Task 2 + 3 + 4 - Principal Component Analysis
- 3 Task 5 - Non-negative matrix factorization**
- 4 K-Means Clustering

Repetition: Non-negative matrix factorization

- For some data PCA is not intuitive
- Example: Non-negative data
 - Principal directions will have negative entries
 - This can be hard to interpret
- Many datasets are strictly positive
 - Text data
 - Image data
 - Probabilistic data
- Very easy to implement (as the PCA algorithm)

From PCA to NMF

- Given data $X \in \mathbb{R}^{d \times n}$ and PCs in the columns of a matrix $W \in \mathbb{R}^{d \times k}$
- Data can be projected onto PCs $H = W^\top X \Leftrightarrow X \approx WH$

From PCA to NMF

- Given data $X \in \mathbb{R}^{d \times n}$ and PCs in the columns of a matrix $W \in \mathbb{R}^{d \times k}$
- Data can be projected onto PCs $H = W^\top X \Leftrightarrow X \approx WH$
- If $X \geq 0$, it should be possible to choose W, H such that $W \geq 0, H \geq 0$

From PCA to NMF

- Given data $X \in \mathbb{R}^{d \times n}$ and PCs in the columns of a matrix $W \in \mathbb{R}^{d \times k}$
- Data can be projected onto PCs $H = W^\top X \Leftrightarrow X \approx WH$
- If $X \geq 0$, it should be possible to choose W, H such that $W \geq 0, H \geq 0$
- Goal of NMF is to find $W \in \mathbb{R}^{d \times k}$ and $H \in \mathbb{R}^{k \times n}$, such that $\|X - WH\|_{\text{Fro}}^2$ is minimized

From PCA to NMF

- Given data $X \in \mathbb{R}^{d \times n}$ and PCs in the columns of a matrix $W \in \mathbb{R}^{d \times k}$
- Data can be projected onto PCs $H = W^\top X \Leftrightarrow X \approx WH$
- If $X \geq 0$, it should be possible to choose W, H such that $W \geq 0, H \geq 0$
- Goal of NMF is to find $W \in \mathbb{R}^{d \times k}$ and $H \in \mathbb{R}^{k \times n}$, such that $\|X - WH\|_{\text{Fro}}^2$ is minimized
- Cannot be directly calculated \Rightarrow iterative optimization with gradient descent

NMF Algorithm

Algorithm 3: Non-negative Matrix Factorization

Input: Dataset $X \in \mathbb{R}_+^{d \times n}$; Dimensionality k

Output: $W \in \mathbb{R}_+^{d \times k}$, Hidden causes $H \in \mathbb{R}_+^{k \times n}$ of the dataset

- 1 Initialize $W \in \mathbb{R}_+^{d \times k}$ and $H \in \mathbb{R}_+^{k \times n}$ randomly
 - 2 Add a small constant $\sigma = 10^{-19}$ to avoid zero-divisions
 - 3 **for** $it \leq iterations$ **do**
 - 4 $H \leftarrow H + \eta (W^\top W H - X^\top W)$
 - 5 $W \leftarrow W + \eta (W H H^\top - X H^\top)$
 - 6 **return** W and H
-

Task 5 - NMF (actually simple linear algebra)

- NMF applied to a dataset $X \in \mathbb{R}^{4 \times 3}$
- After training the reconstruction $\tilde{X} = WH$ looks like this

$$\tilde{X} = \begin{pmatrix} 1 & 0 & 1 \\ 1 & 1 & 2 \\ 0 & 0 & 2 \\ 0 & 1 & 1 \end{pmatrix}, \quad H = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- What is W ?

Task 5 - NMF (actually simple linear algebra)

- NMF applied to a dataset $X \in \mathbb{R}^{4 \times 3}$
- After training the reconstruction $\tilde{X} = WH$ looks like this

$$\tilde{X} = \begin{pmatrix} 1 & 0 & 1 \\ 1 & 1 & 2 \\ 0 & 0 & 2 \\ 0 & 1 & 1 \end{pmatrix}, \quad H = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

- What is W ?
- $\tilde{X} = WH = W$

- 1 Task 1 + 4 - Eigenvalues and Eigenvectors
- 2 Task 2 + 3 + 4 - Principal Component Analysis
- 3 Task 5 - Non-negative matrix factorization
- 4 K-Means Clustering**

K-Means Clustering

- Theory is completely covered by the lecture, but there are also Tutorials on youtube
- Explanation of the concept with the help of the solution of the quiz

⇒ Jupyter Notebook